

# Scalable Machine Learning in H2O

July 2018



Navdeep Gill  
@Navdeep\_Gill\_

# Agenda

- H2O Introduction
- H2O Core Overview
- H2O R API
- Demos

# Company Overview

<b>Founded</b>	2011 Venture-backed, Debuted in 2012
<b>Products</b>	<ul style="list-style-type: none"><li>• <b>H<sub>2</sub>O Open Source In-Memory AI Prediction Engine</b></li><li>• Sparkling Water (H<sub>2</sub>O + Spark)</li><li>• H2O4GPU (H2O on GPUs)</li><li>• Enterprise Steam</li><li>• Driverless AI</li></ul>
<b>Mission</b>	Operationalize Data Science & Provide a Platform to Build Beautiful Data Products
<b>Team</b>	<p>75+ employees</p> <ul style="list-style-type: none"><li>• Distributed Systems Engineers doing Machine Learning</li><li>• World-class Visualization Designers</li></ul>
<b>Headquarters</b>	Mountain View, CA



3

# Scientific Advisory Council



## Dr. Trevor Hastie

- PhD in Statistics, Stanford University
- John A. Overdeck Professor of Mathematics, Stanford University
- Co-author, *The Elements of Statistical Learning: Prediction, Inference and Data Mining*
- Co-author, *Generalized Additive Models*
- 108,404 citations (via Google Scholar)



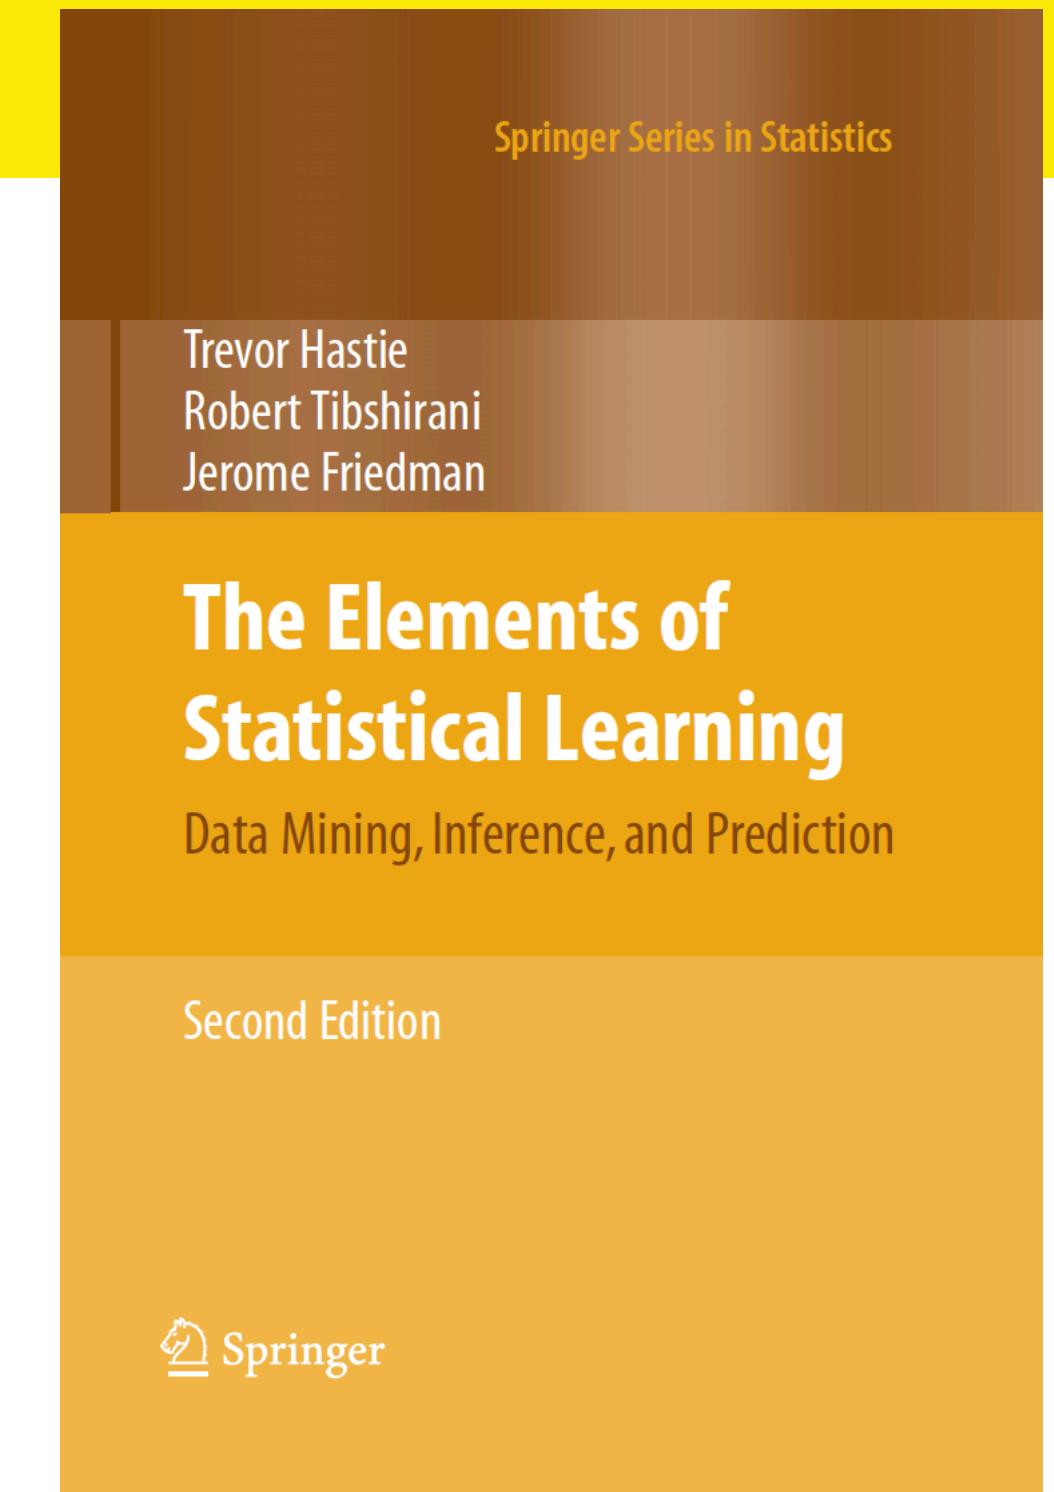
## Dr. Robert Tibshirani

- PhD in Statistics, Stanford University
- Professor of Statistics and Health Research and Policy, Stanford University
- COPPS Presidents' Award recipient
- Co-author, *The Elements of Statistical Learning: Prediction, Inference and Data Mining*
- Author, *Regression Shrinkage and Selection via the Lasso*
- Co-author, *An Introduction to the Bootstrap*



## Dr. Steven Boyd

- PhD in Electrical Engineering and Computer Science, UC Berkeley
- Professor of Electrical Engineering and Computer Science, Stanford University
- Co-author, *Convex Optimization*
- Co-author, *Linear Matrix Inequalities in System and Control Theory*
- Co-author, *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*



# What is H2O?

**Java-Based Software for In-Memory Data Modeling**

**Open Source**



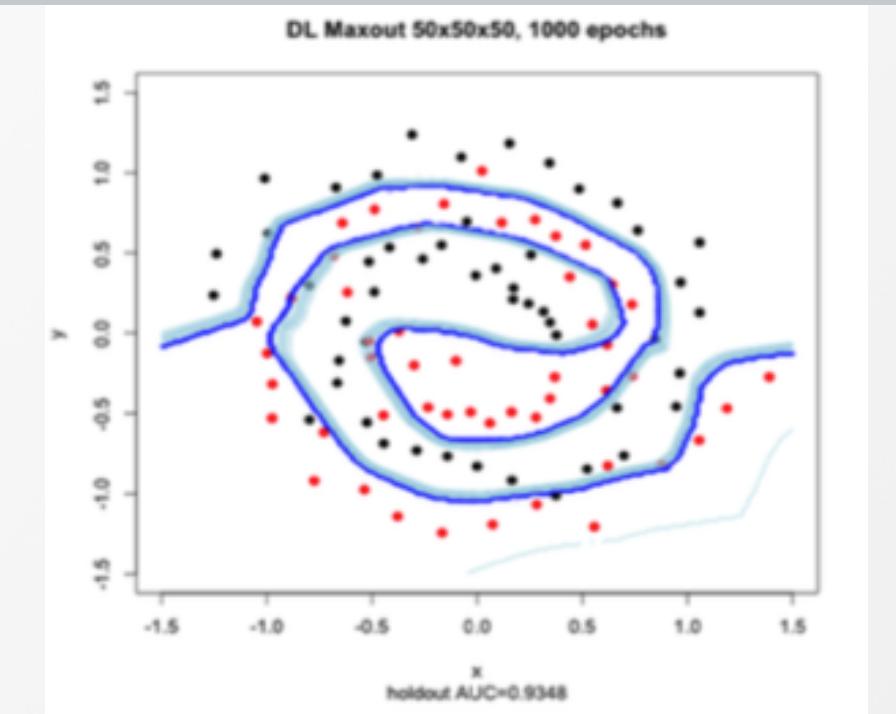
**Big Data Ecosystem**



**Flexible Interface**

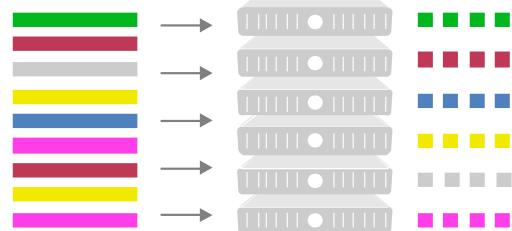
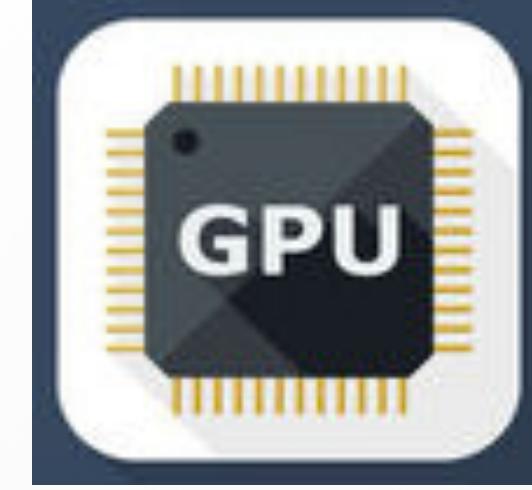


**Smart and Fast Algorithms**

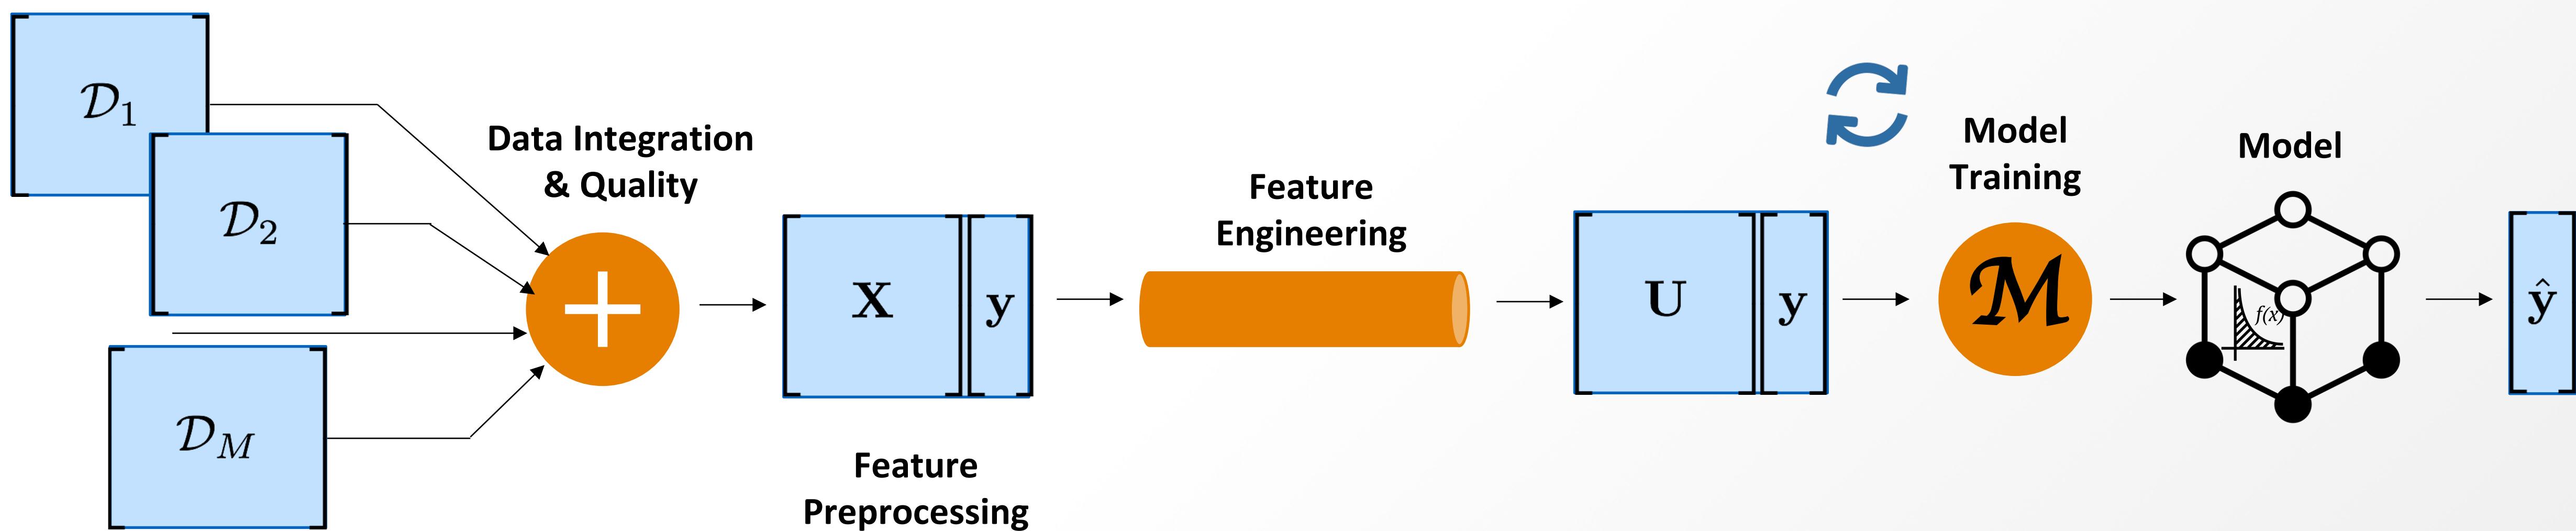


# What is H2O?

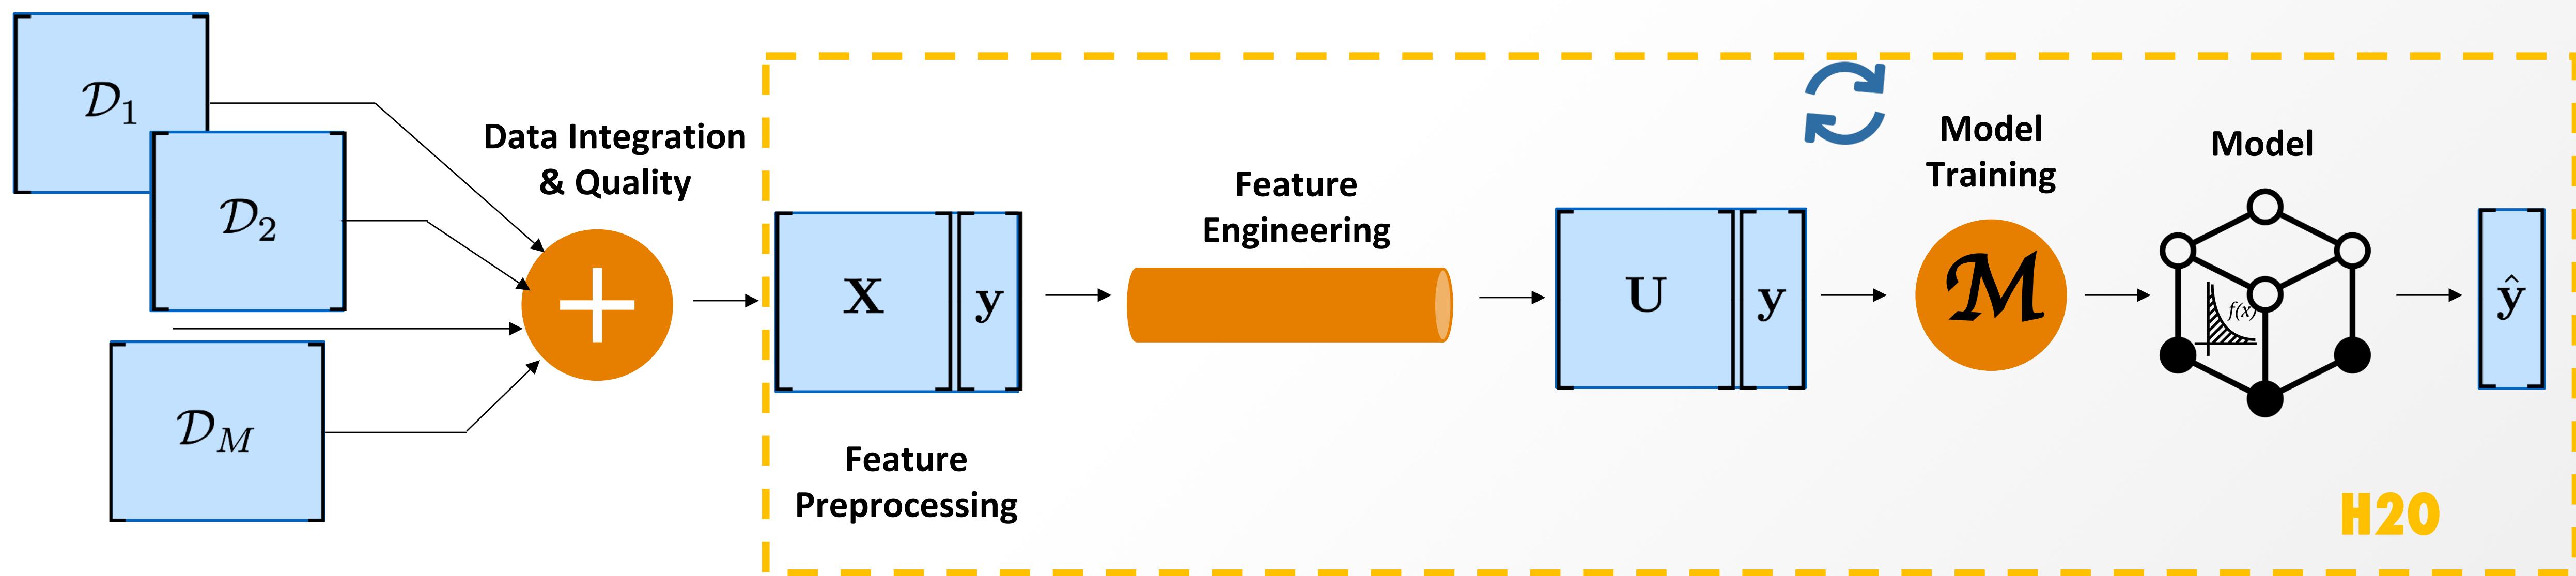
## Java-Based Software for In-Memory Data Modeling

Scalability and Performance	Rapid Model Deployment	GPU Enablement*	Cloud Integration
 <ul style="list-style-type: none"><li>Distributed In-Memory Computing Platform</li><li>Distributed Algorithms</li><li>Fine-Grain MapReduce</li></ul>	<ul style="list-style-type: none"><li>Highly portable models deployed in Java (POJO)</li><li>Automated and streamlined scoring service deployment with Rest API*</li></ul>		  

# The Machine Learning Pipeline



# Where H2O Fits



# Current Algorithm Overview

## Statistical Analysis

---

- Linear Models (GLM)
- Naïve Bayes

## Ensembles

---

- Random Forest
- Gradient Boosting Machine
- Stacking / Super Learner

## Deep Neural Networks

---

- MLP
- Autoencoder
  - Anomaly Detection
  - Deep Features

## Clustering

---

- K-Means (Auto-K)

## Dimension Reduction

---

- Principal Component Analysis
- Generalized Low Rank Models

## Word Embedding

---

- Word2Vec

## Time Series

---

- iSAX

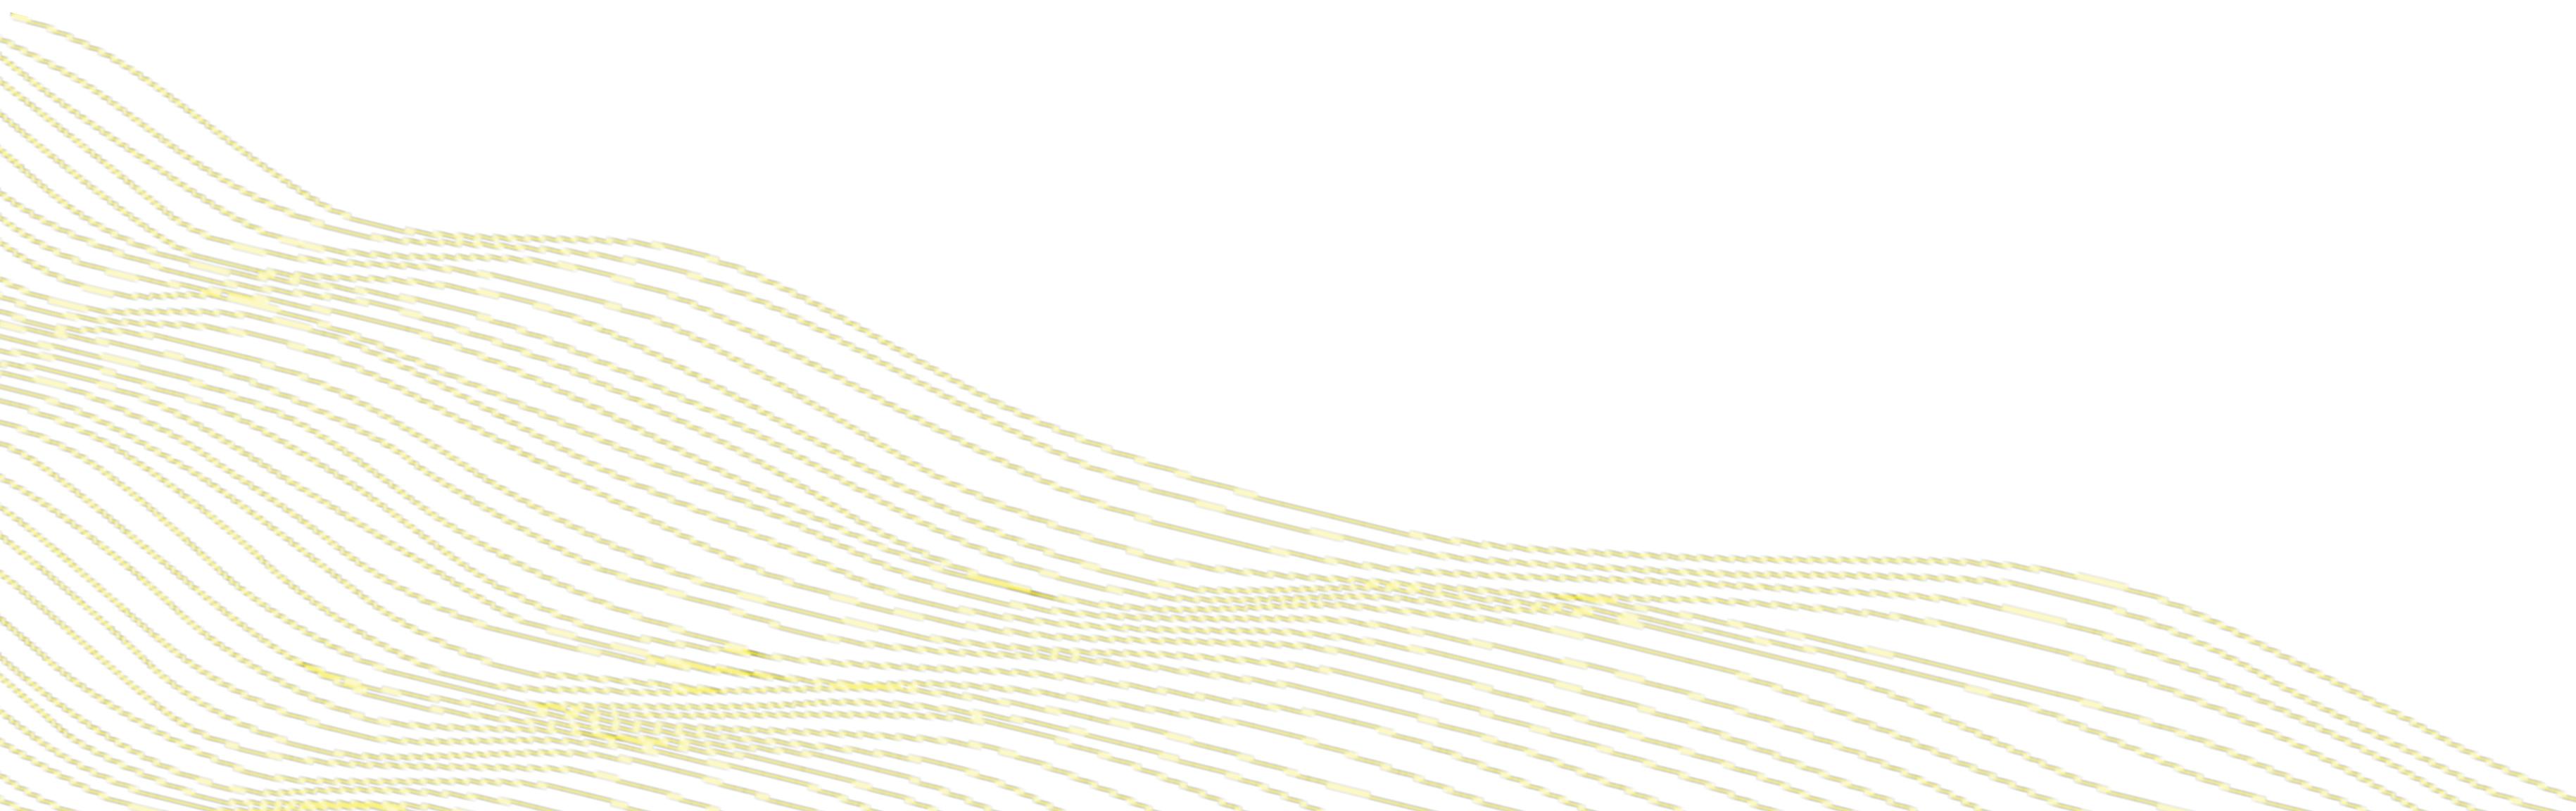
## Machine Learning Tuning

---

- Hyperparameter Search
- Early Stopping

# H2O Core

*Behind the Scenes*



# How H<sub>2</sub>O Core Works



# How H2O Core Works

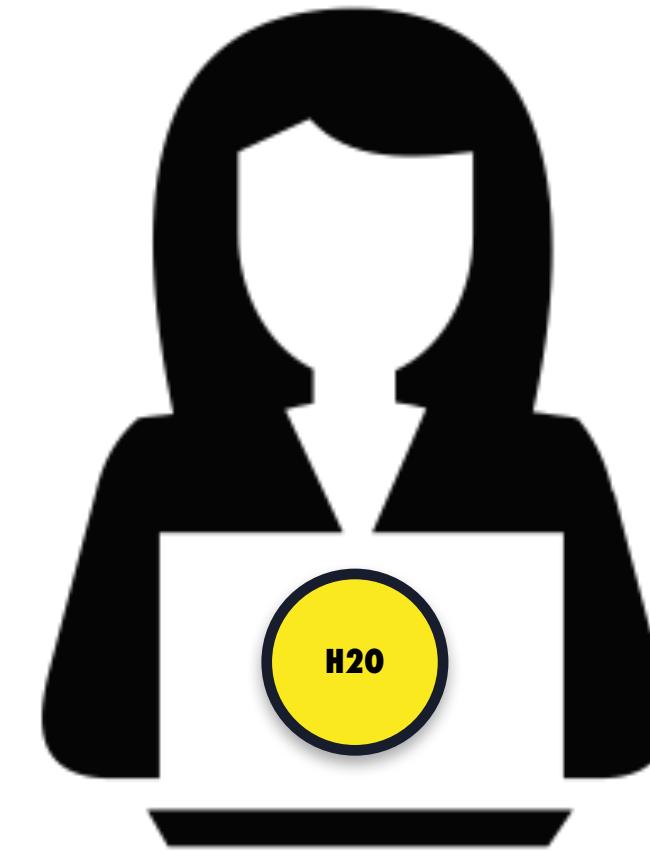


(just a java application)

# How H2O Core Works



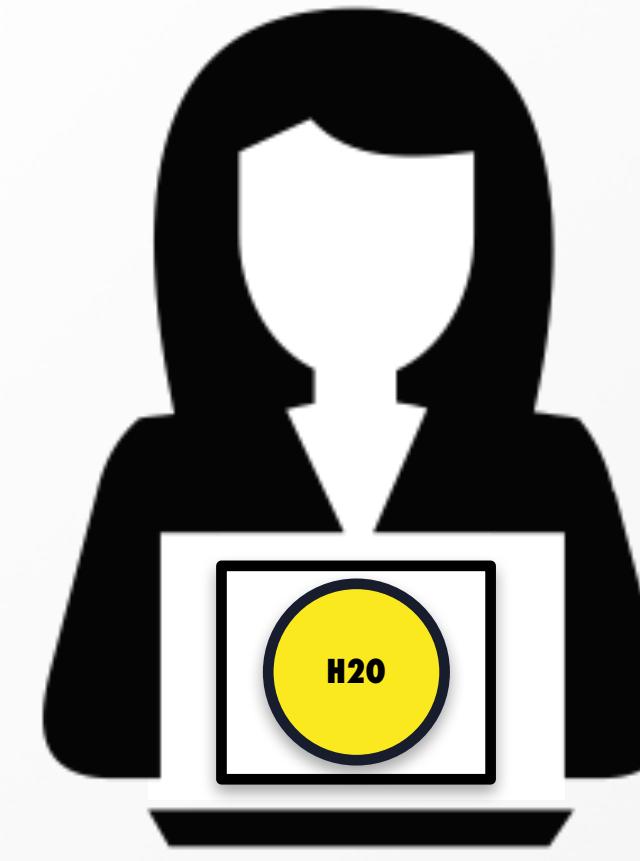
# How H2O Core Works



on a laptop

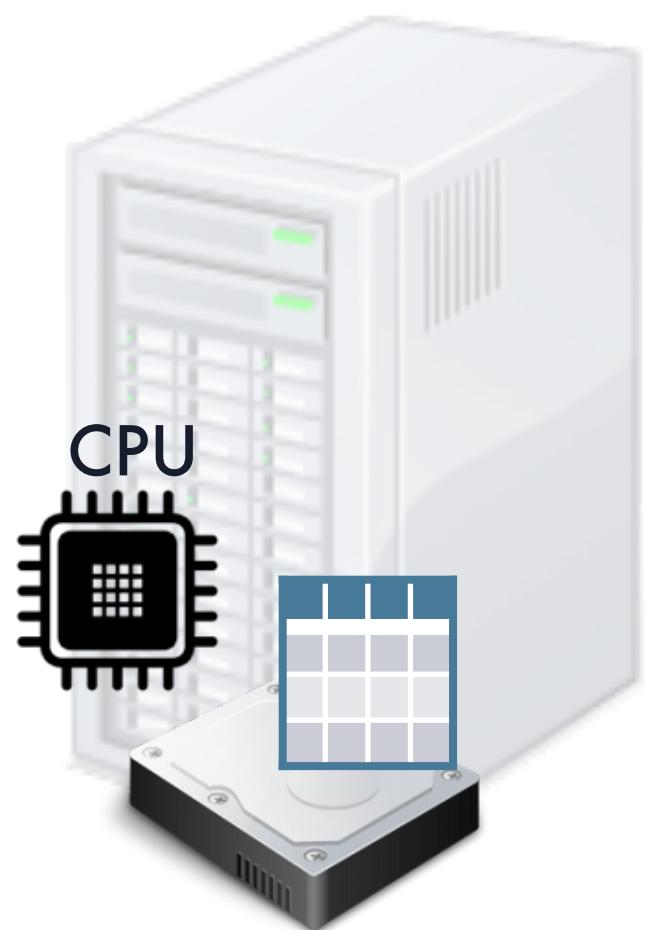


on a virtual machine

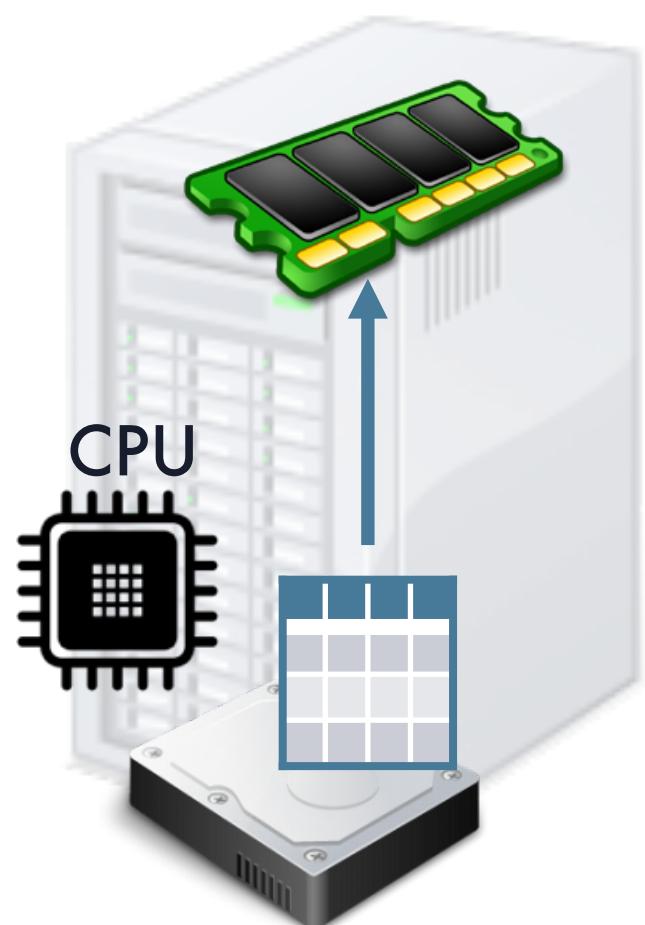


in a container

# H2O Core

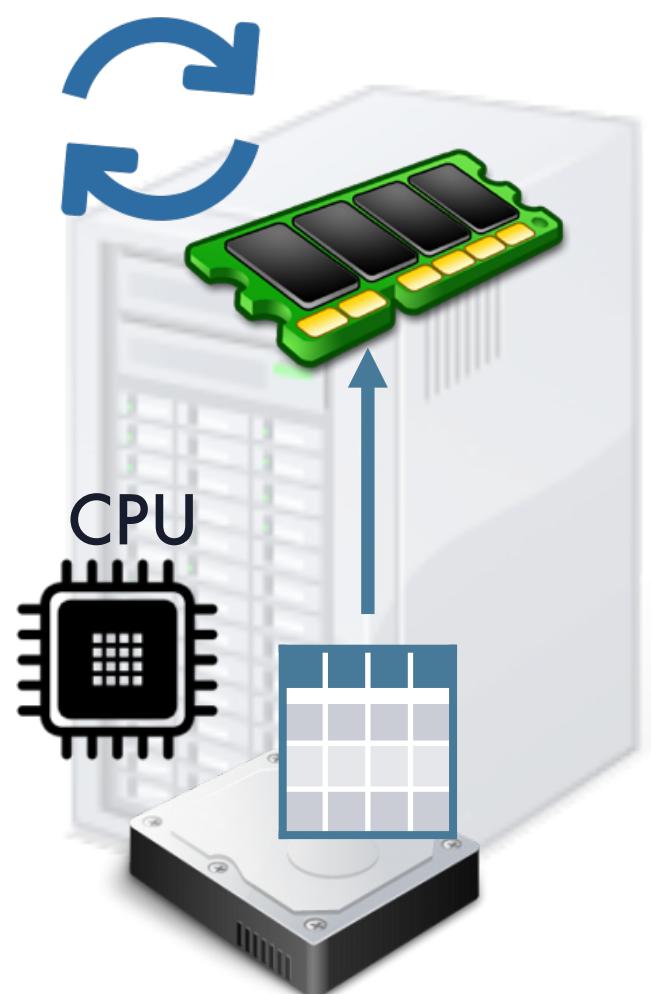


# H2O Core



# H2O Core

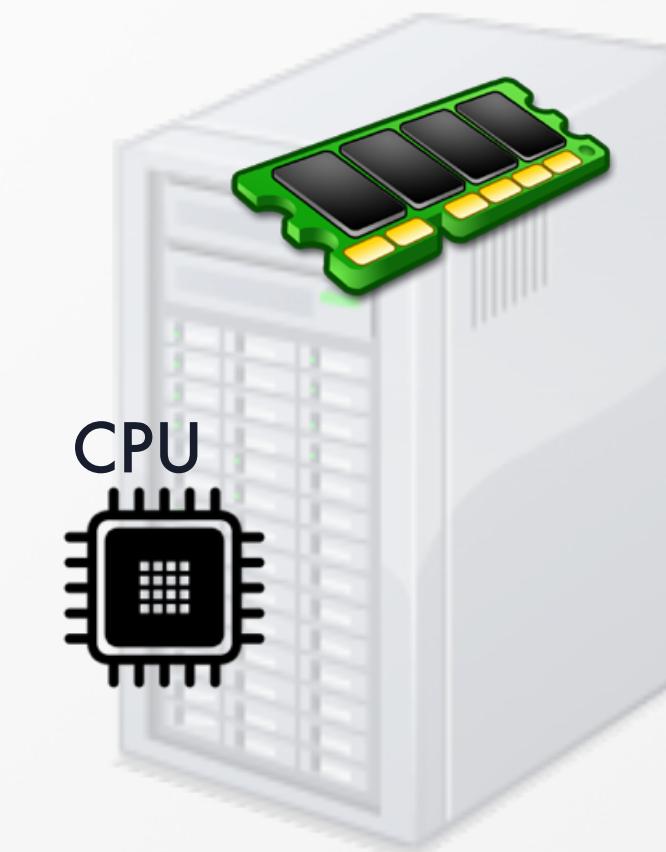
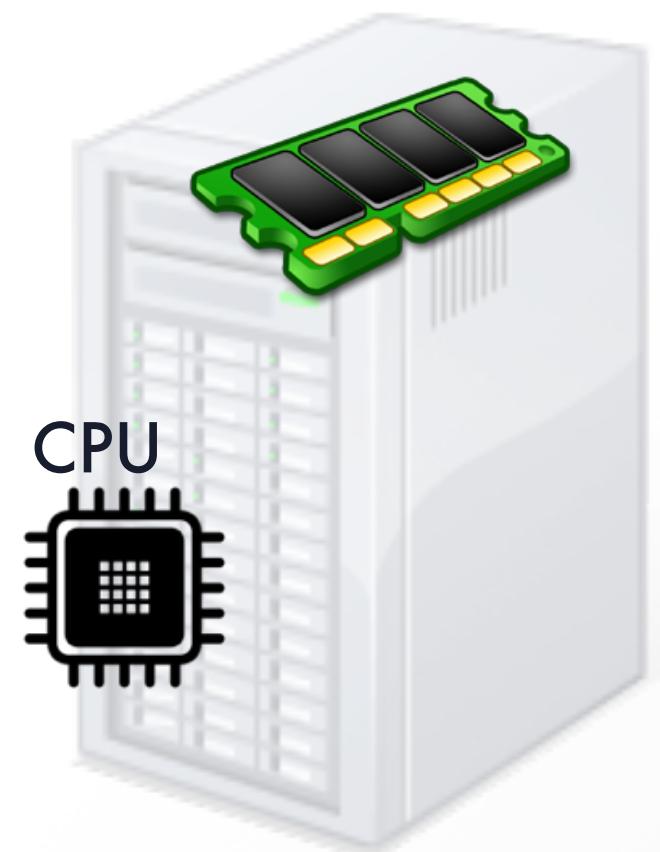
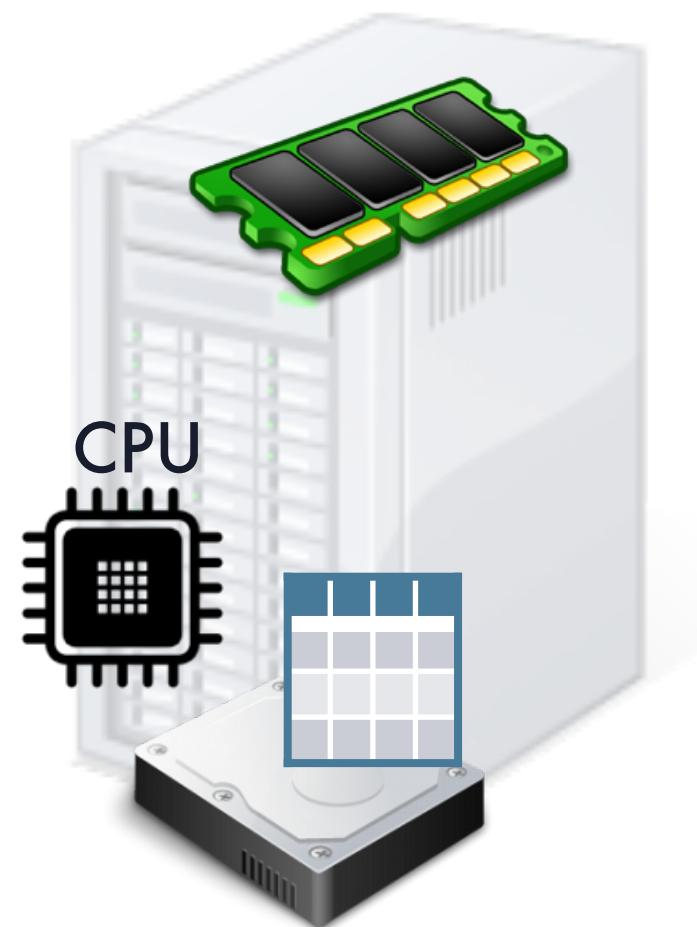
Model Building



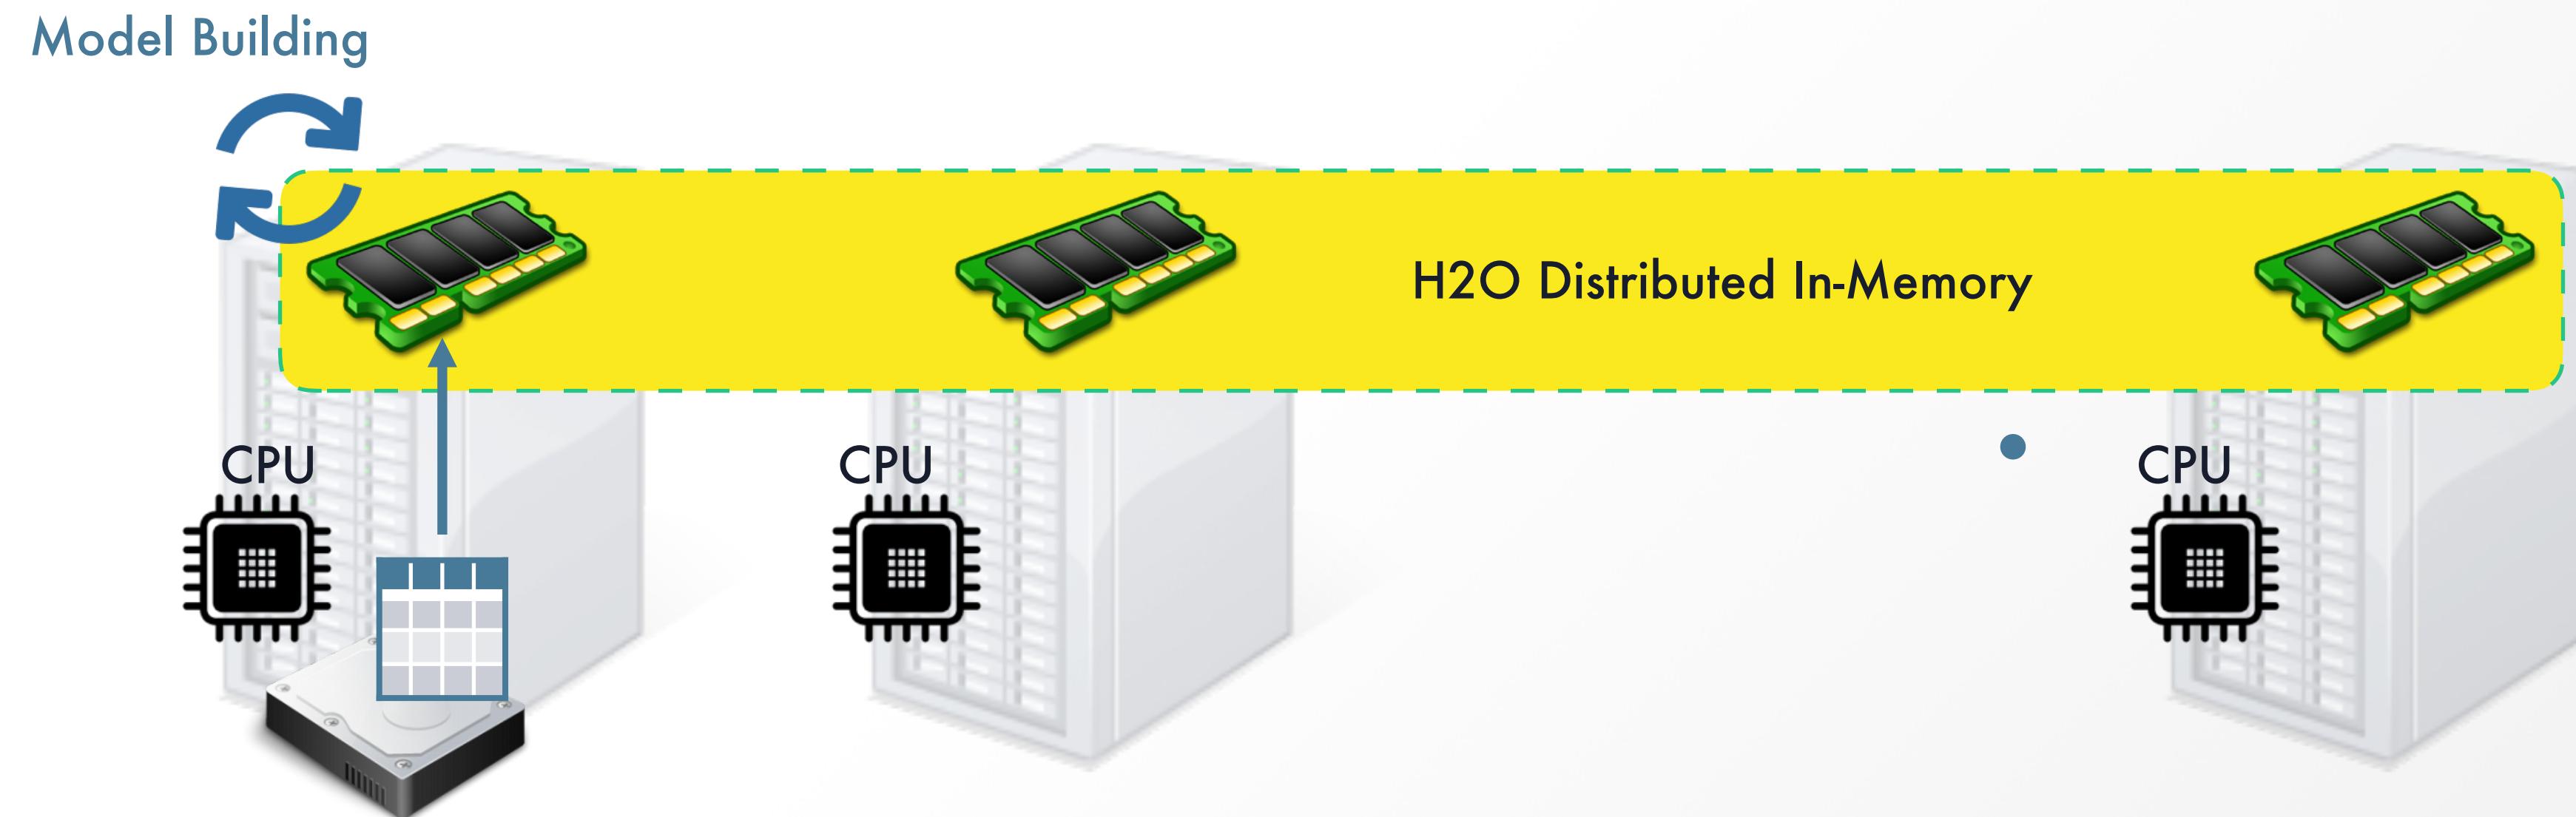
# H2O Core



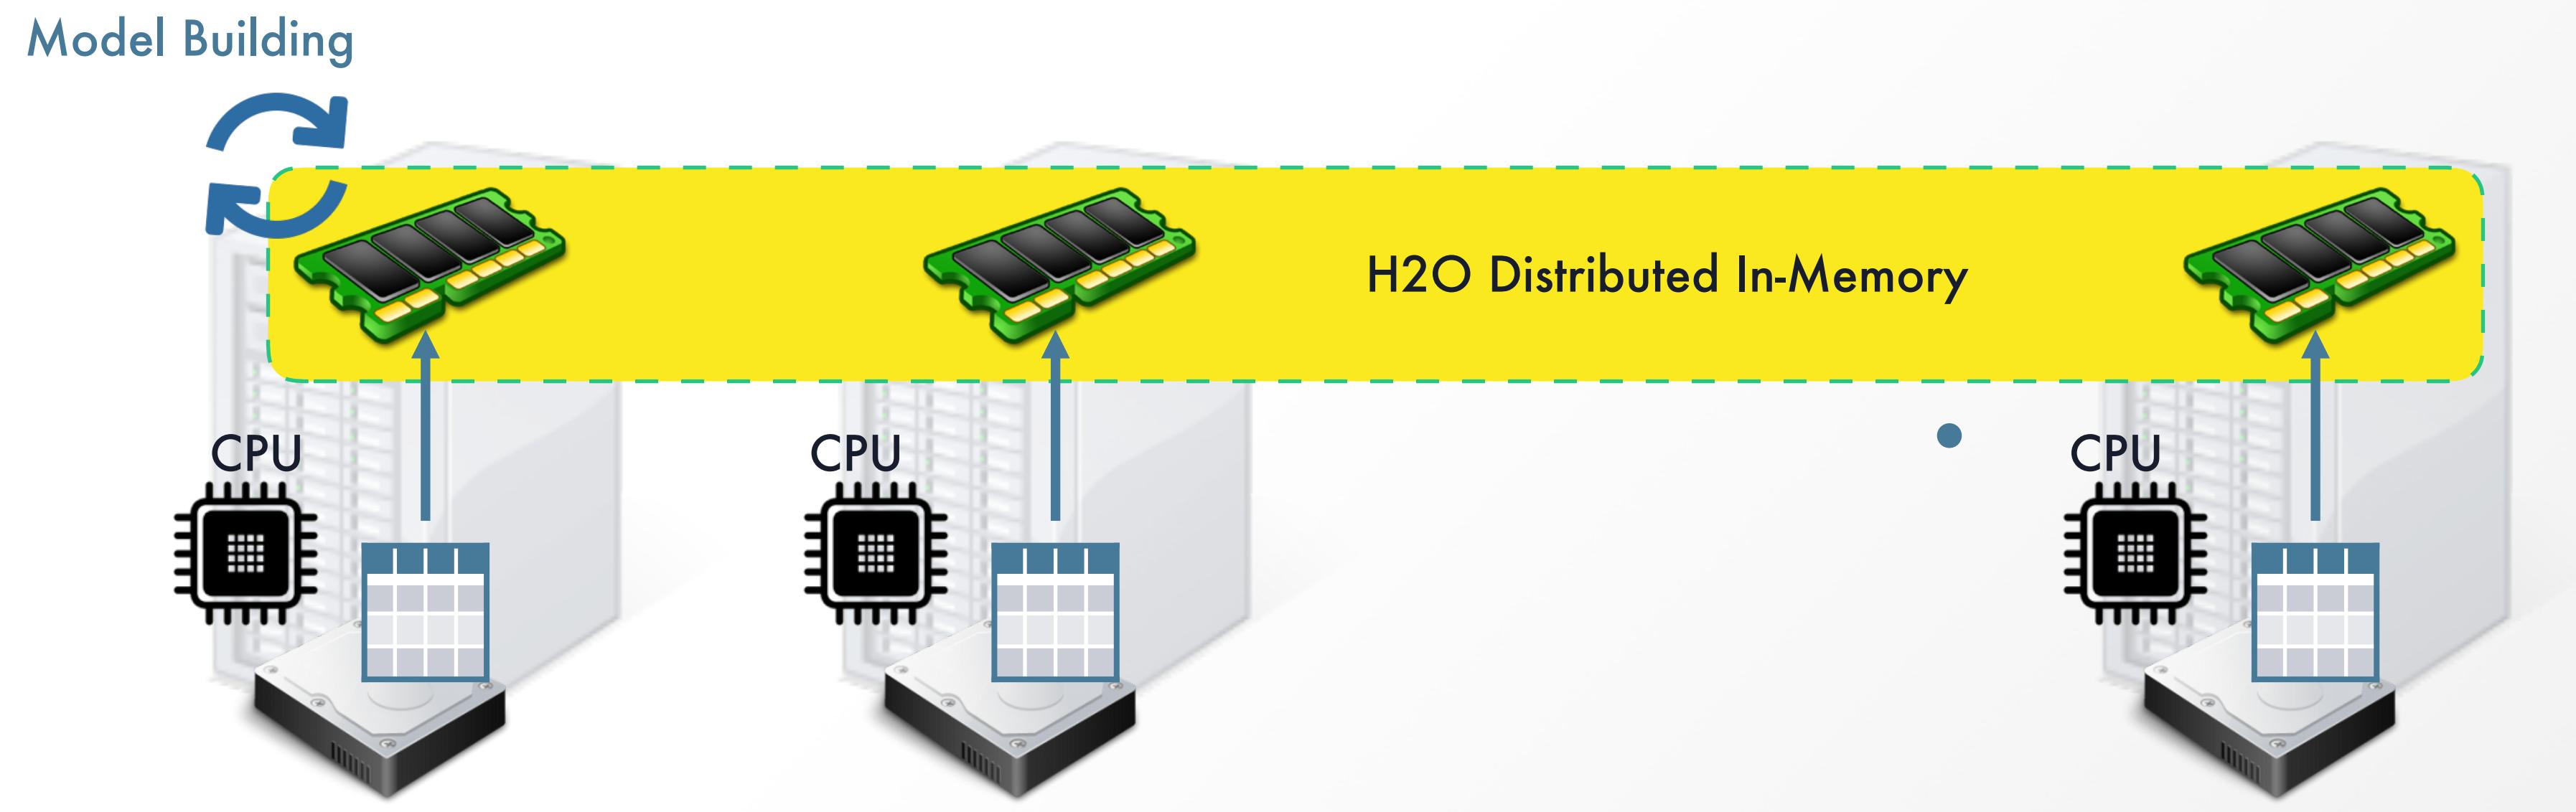
# H2O Core



# H2O Core



# H2O Core with Hadoop



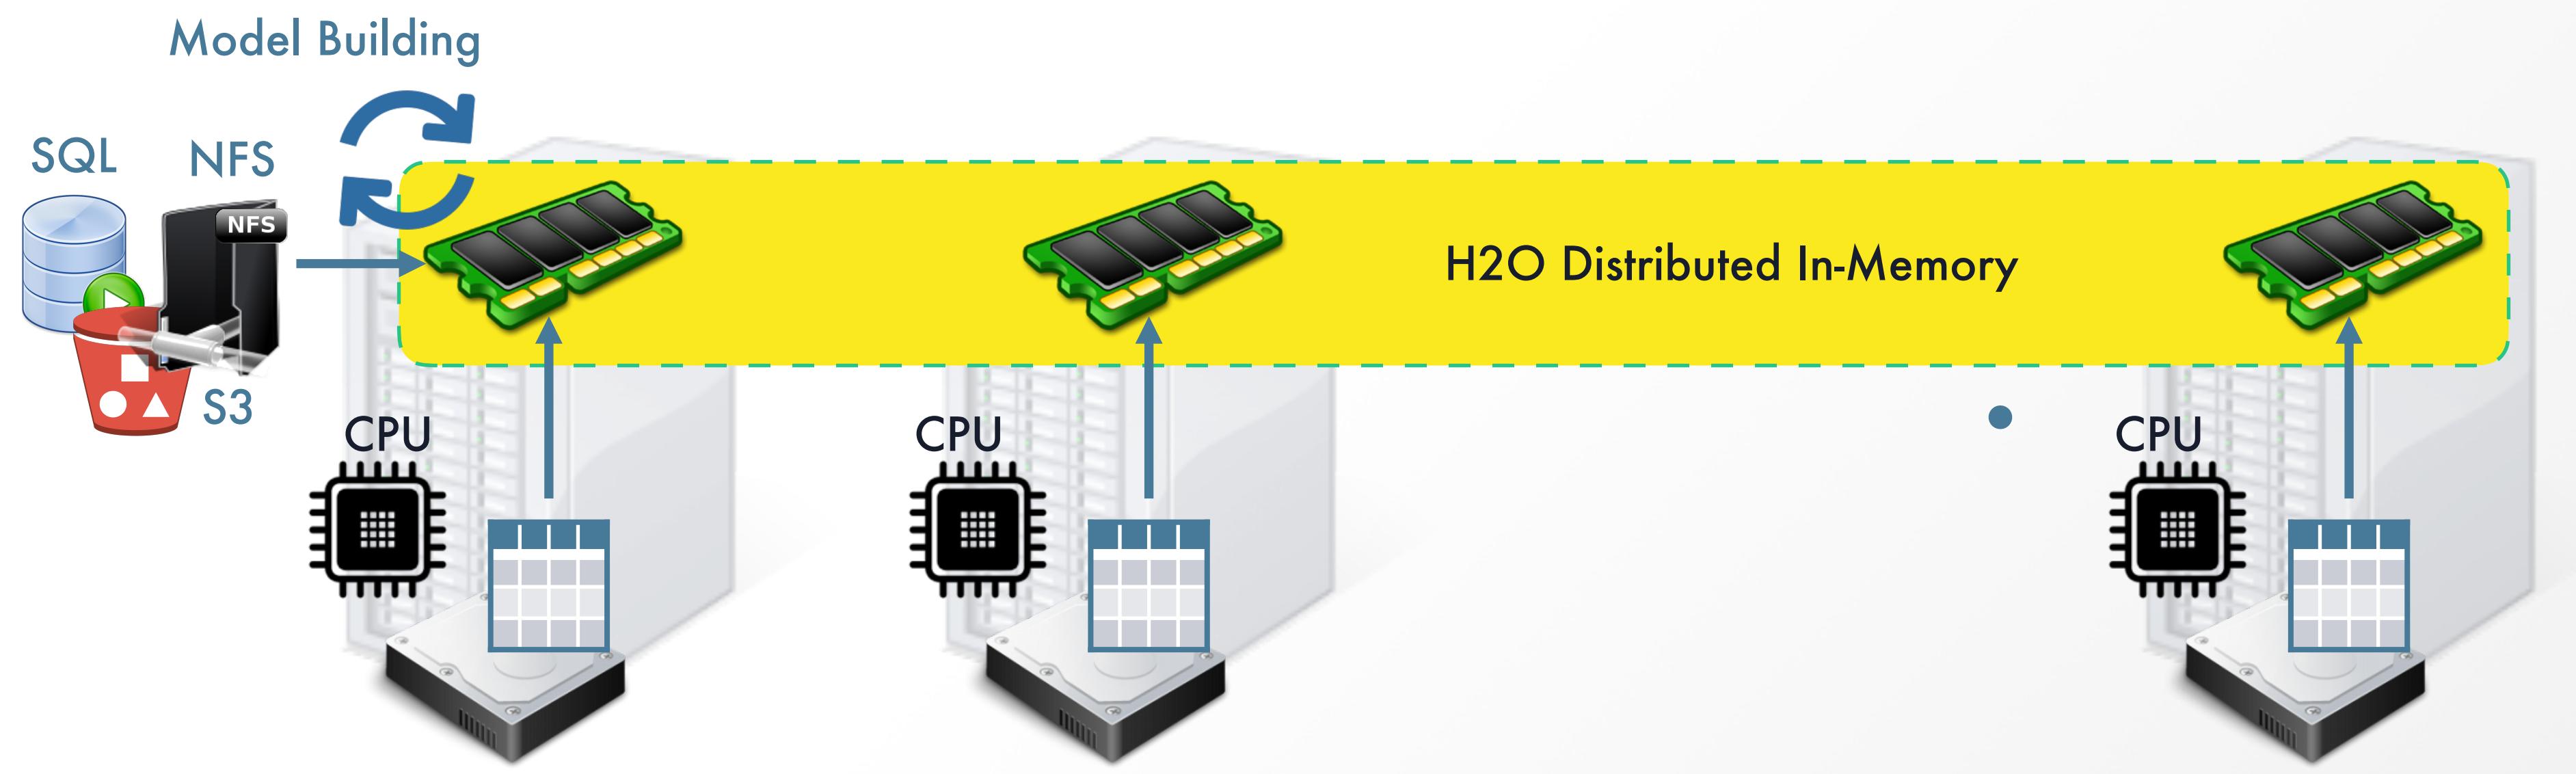
YARN

cloudera Hortonworks

MAPR

H<sub>2</sub>O  
WORLD

# H2O Core with Other Data Sources



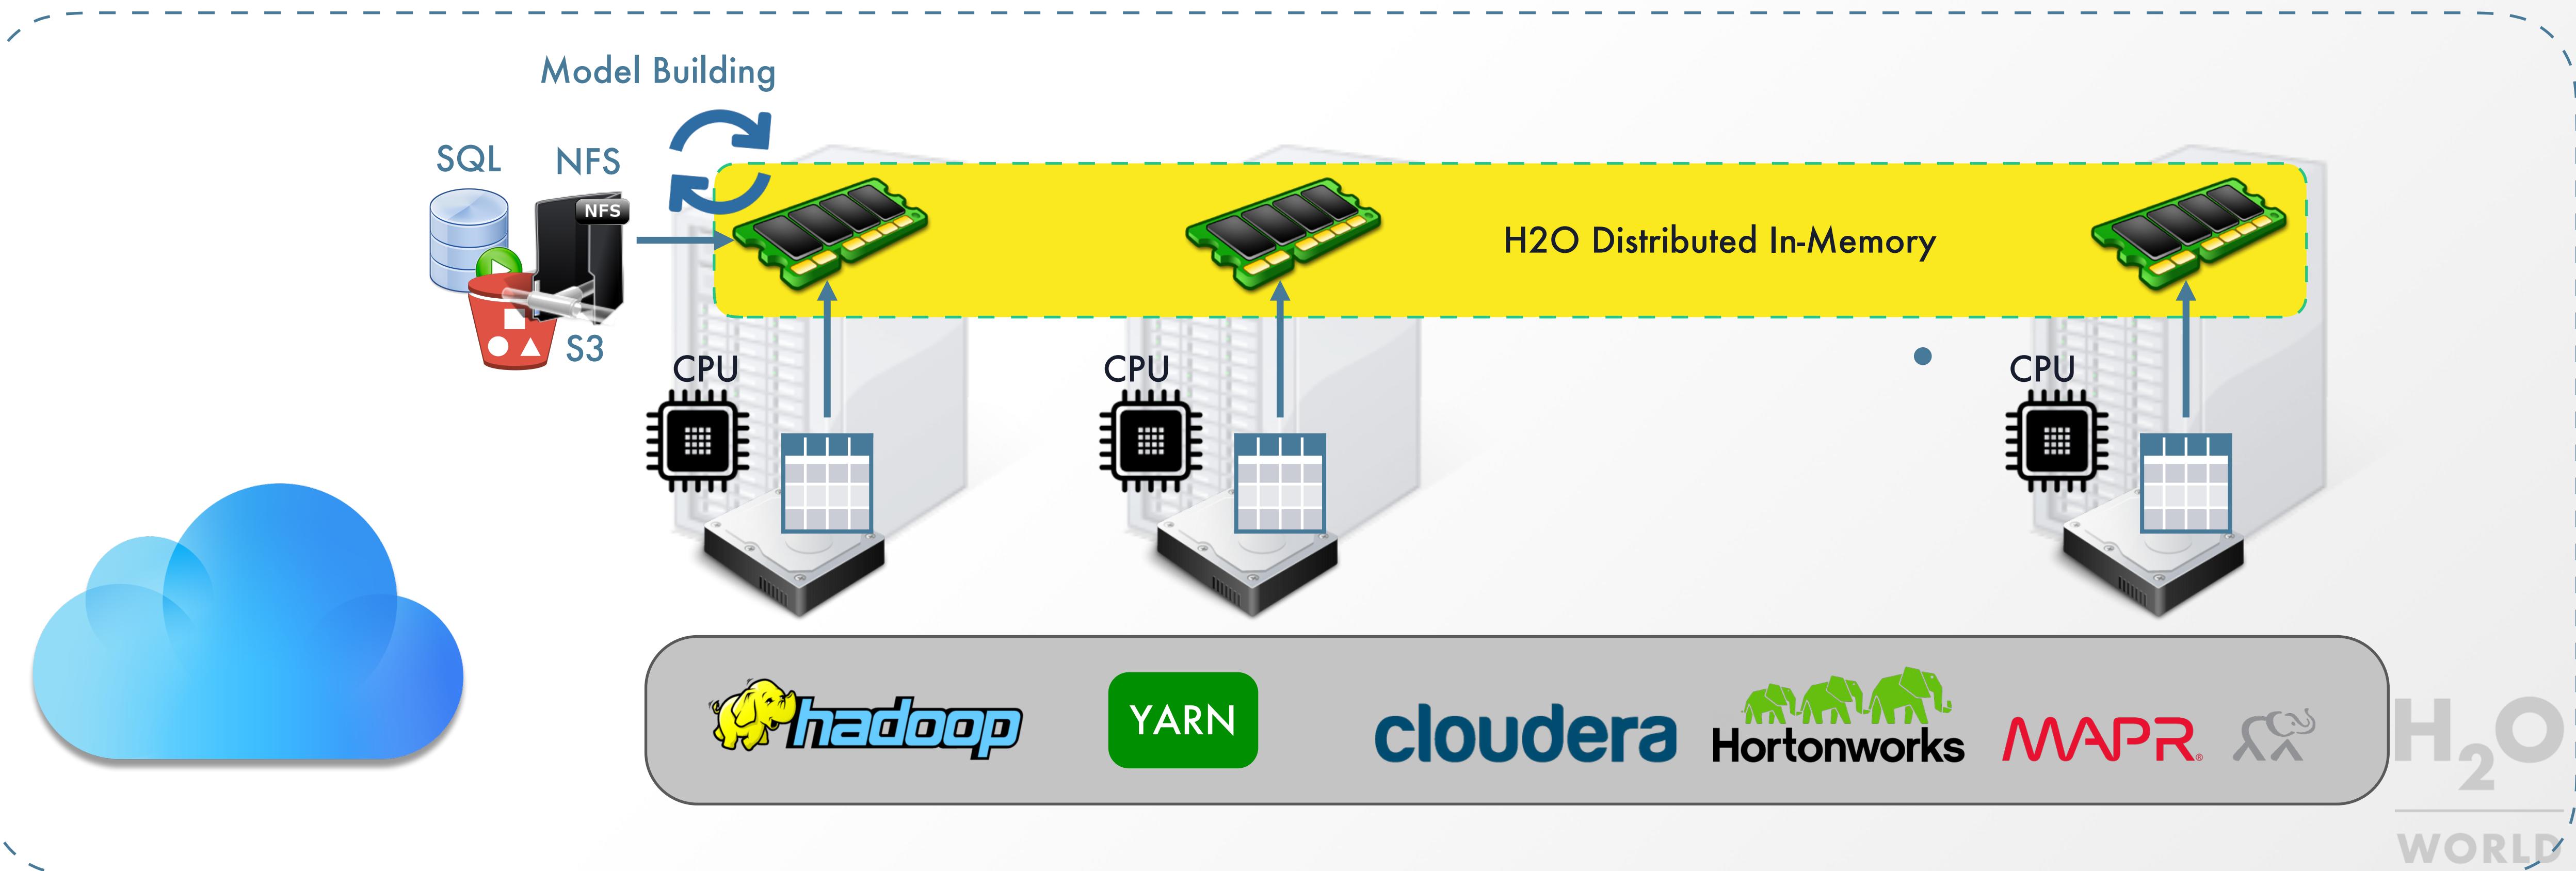
YARN

cloudera Hortonworks

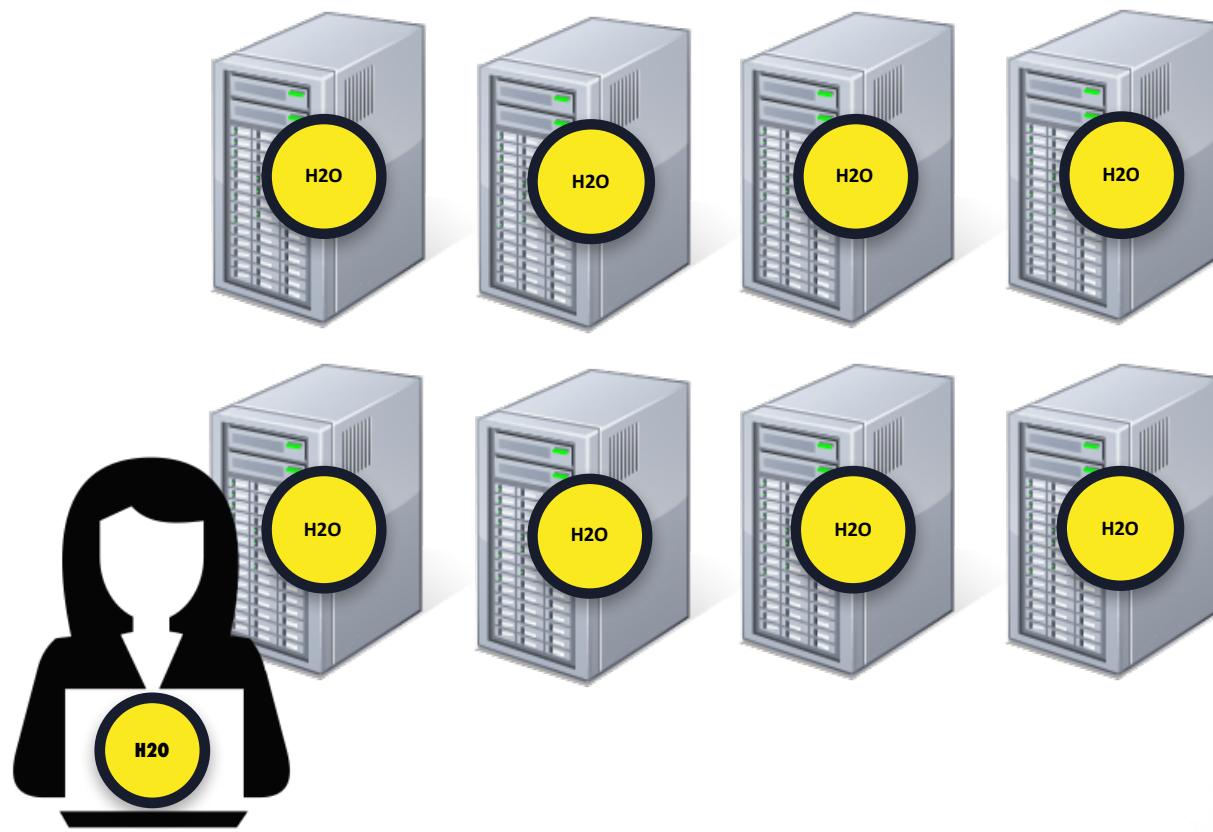
MAPR

H<sub>2</sub>O  
WORLD

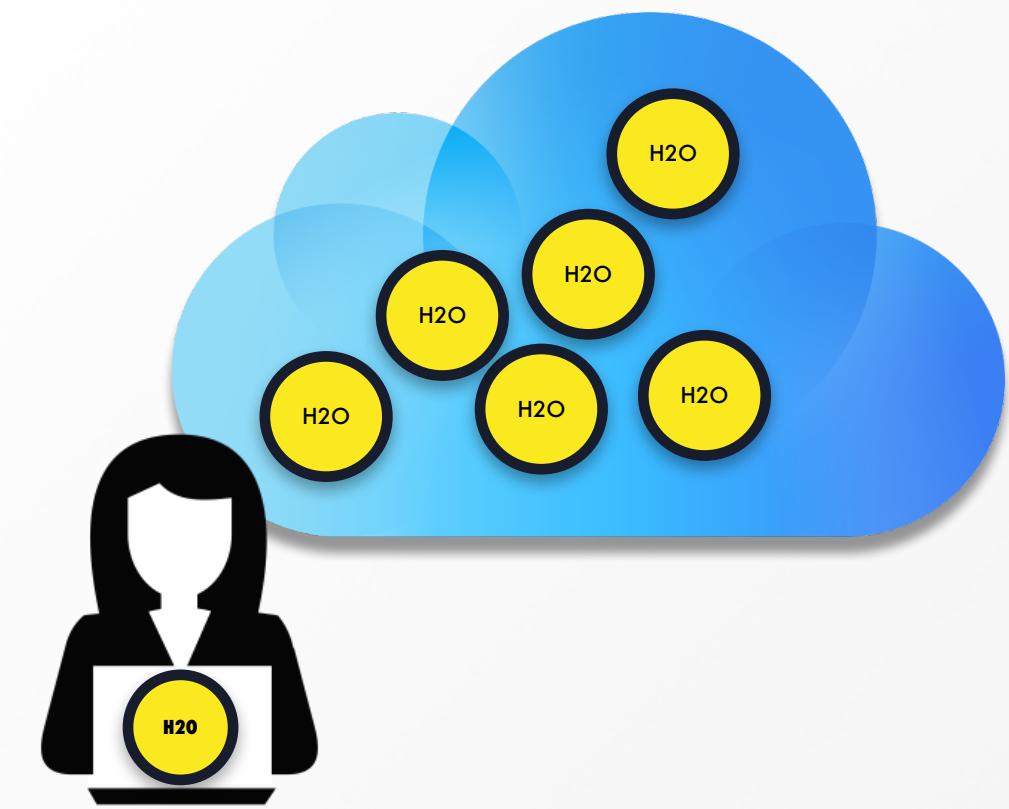
# H2O Core on the Cloud



# H2O Distributed Environments



distributed across multiple machines



distributed on the cloud

# The H2O R API

*(how the client & cluster communicate)*



# H2O R API



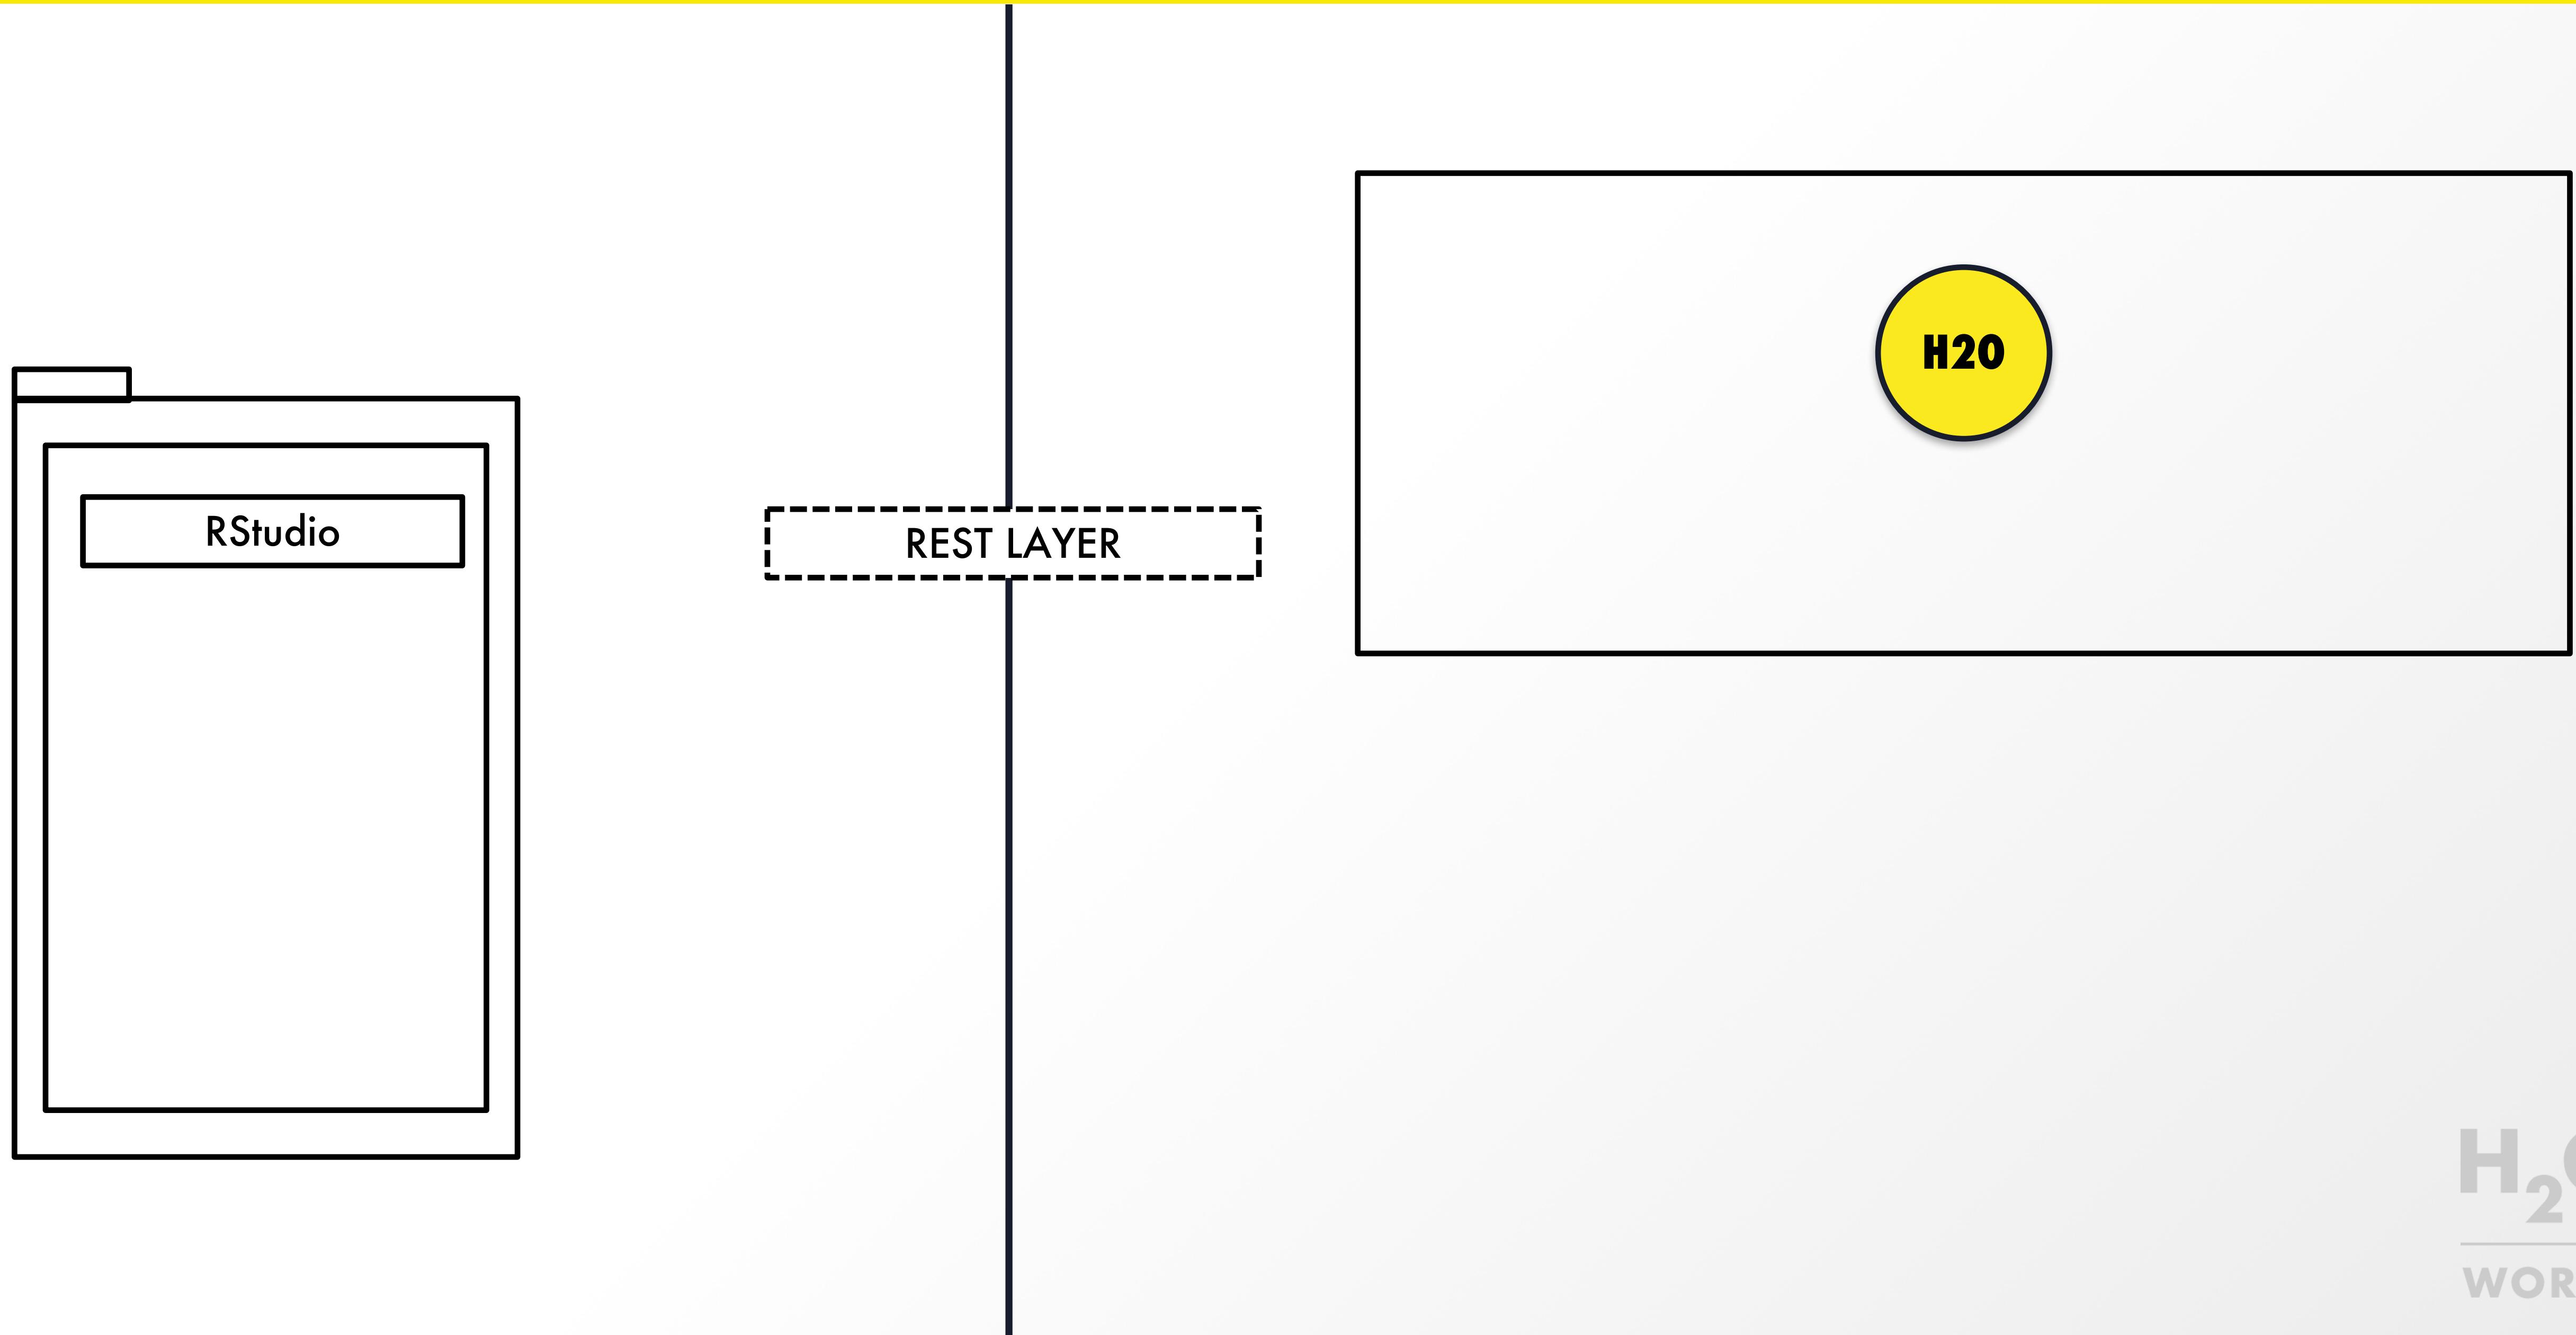
(ALL GREAT THINGS)

# H2O R API

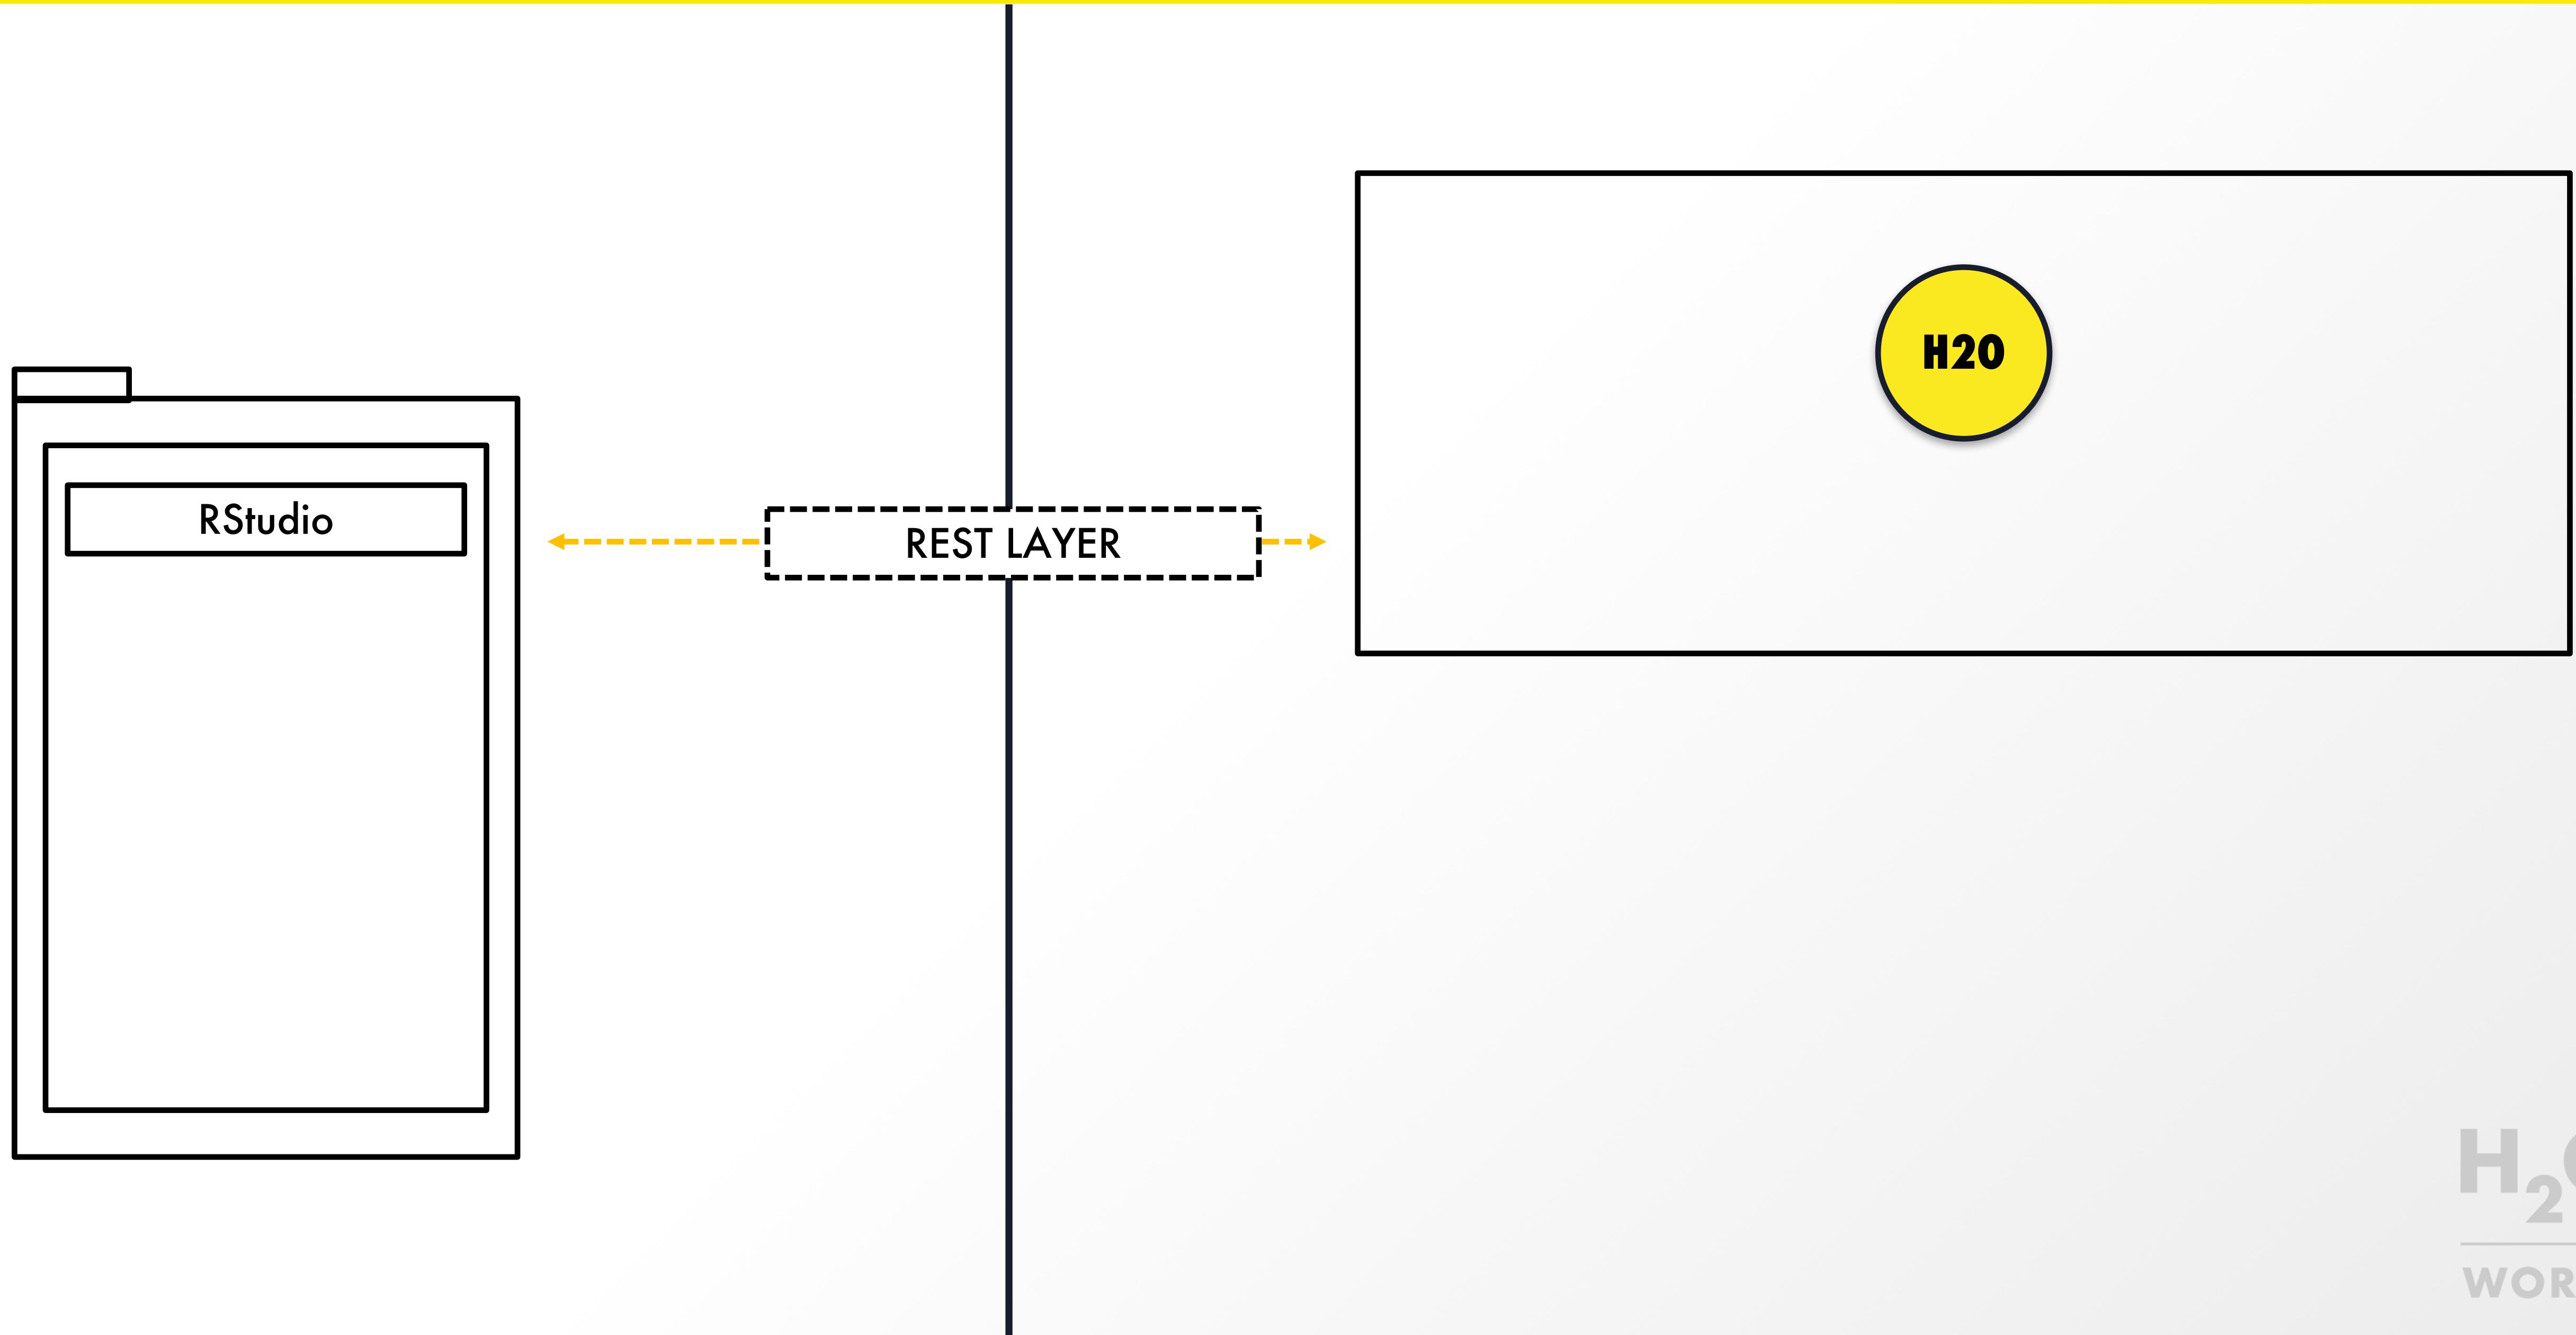


**parallelism & distribution of work**

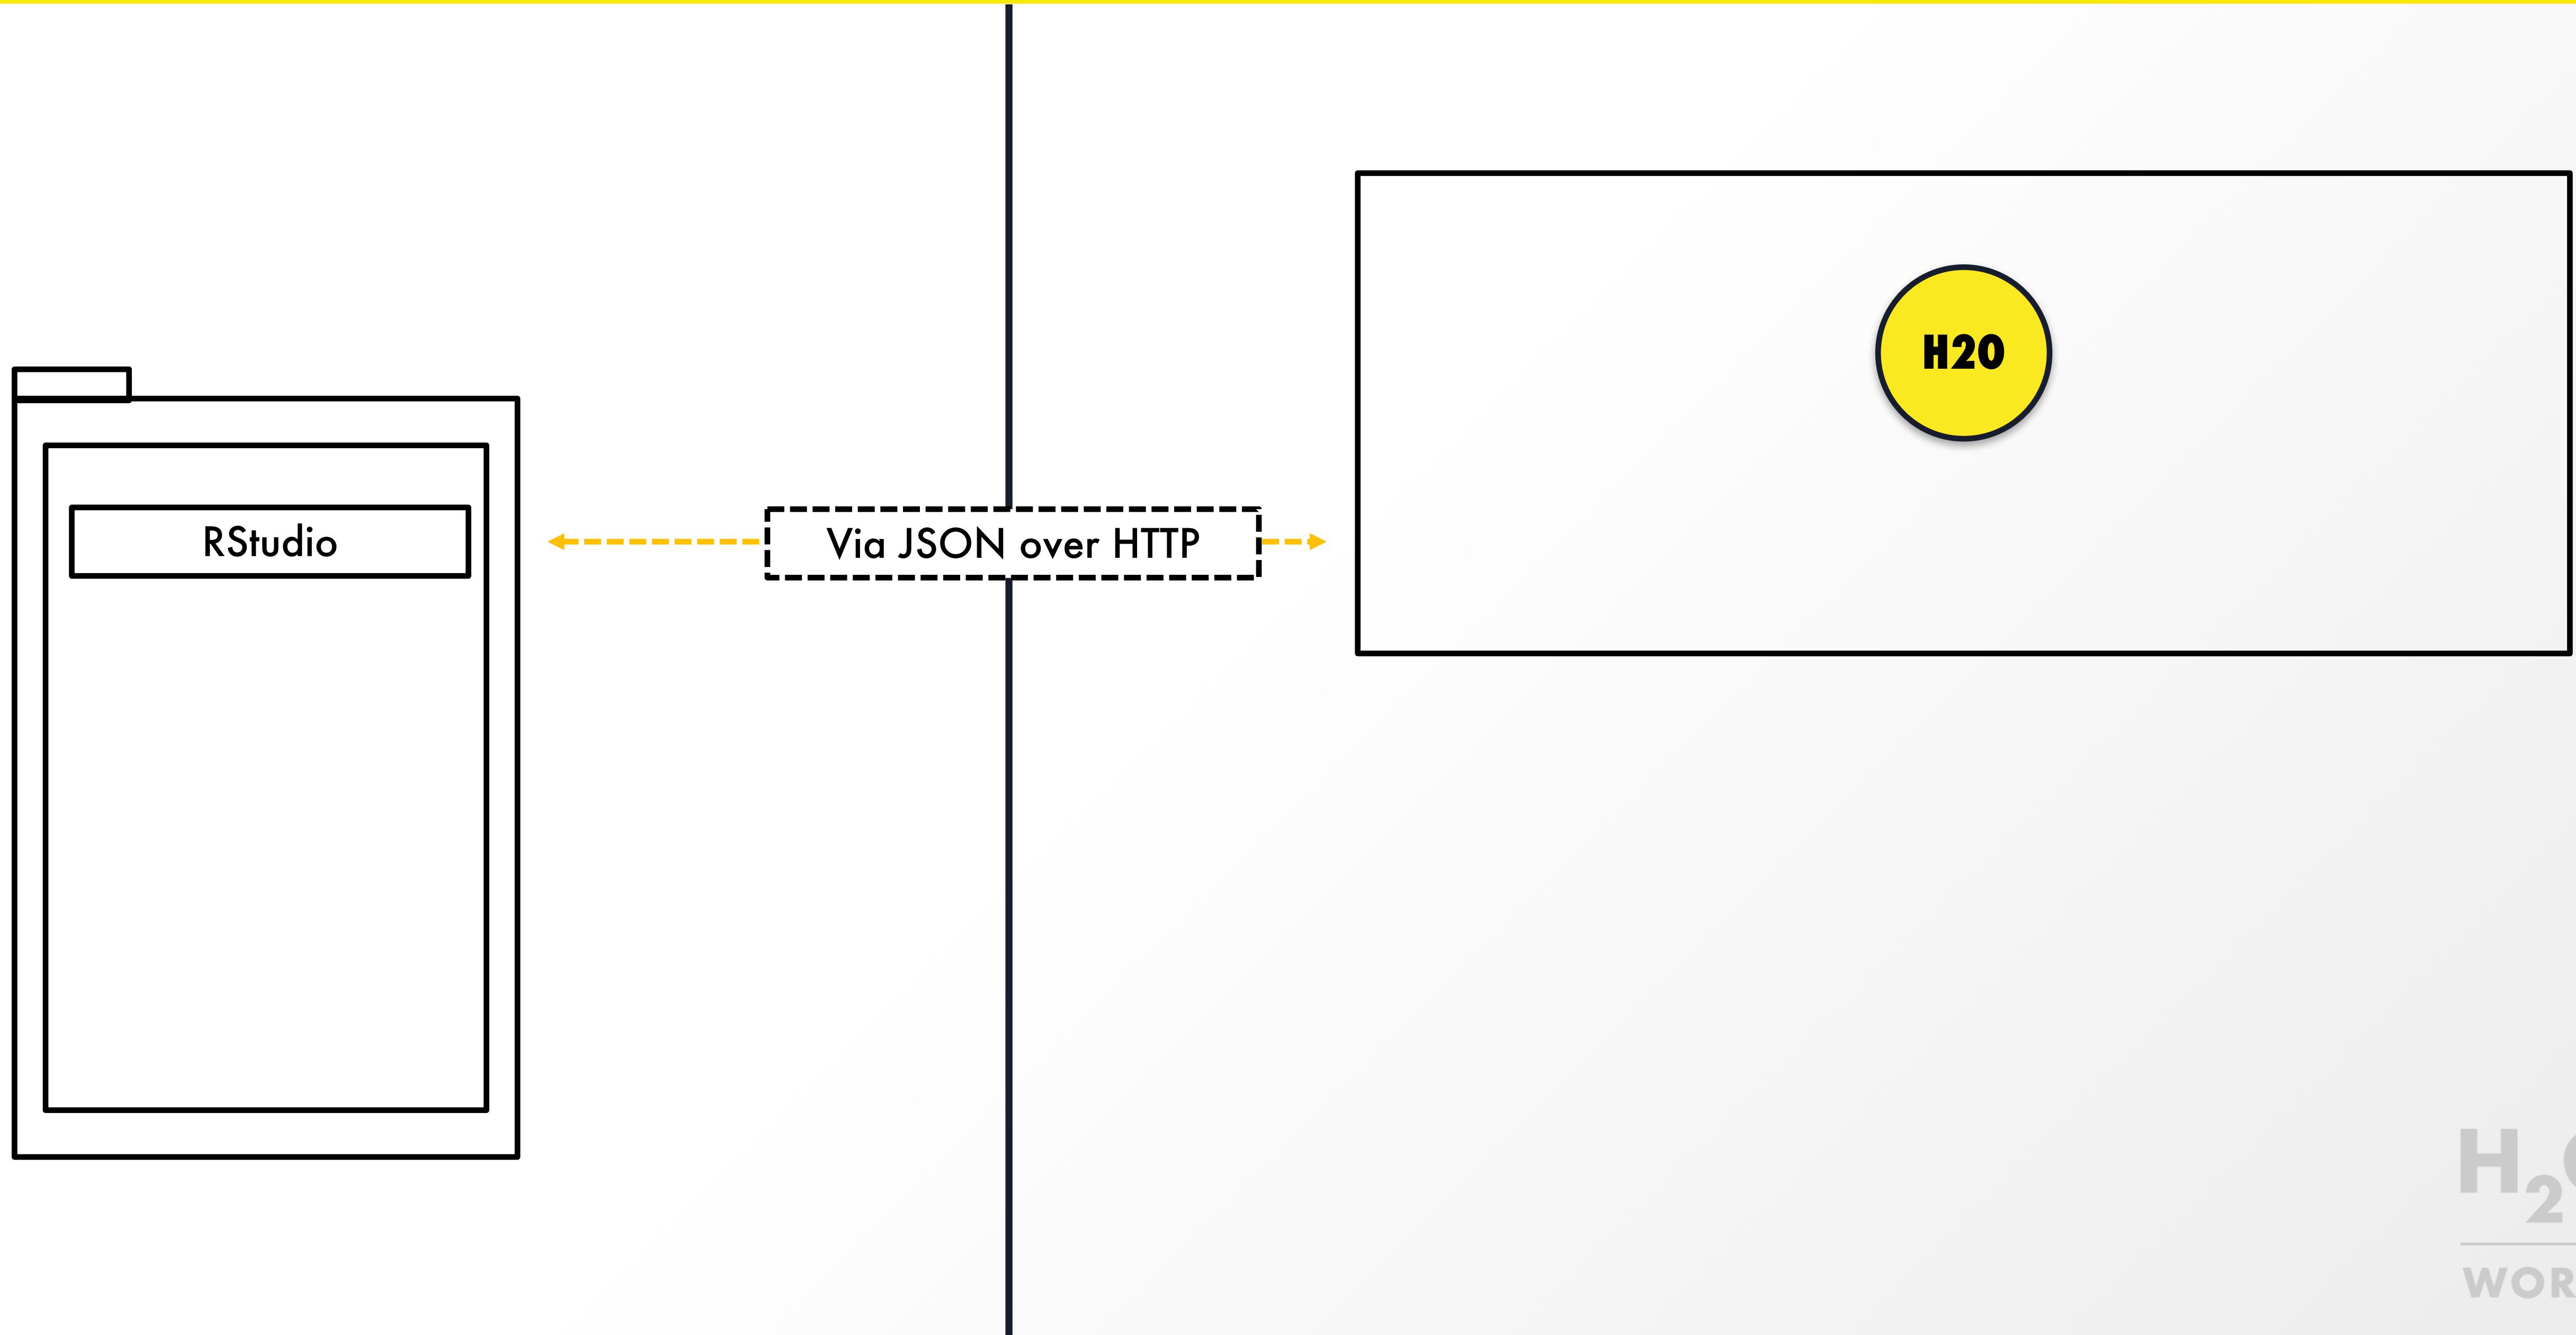
# H2O R API



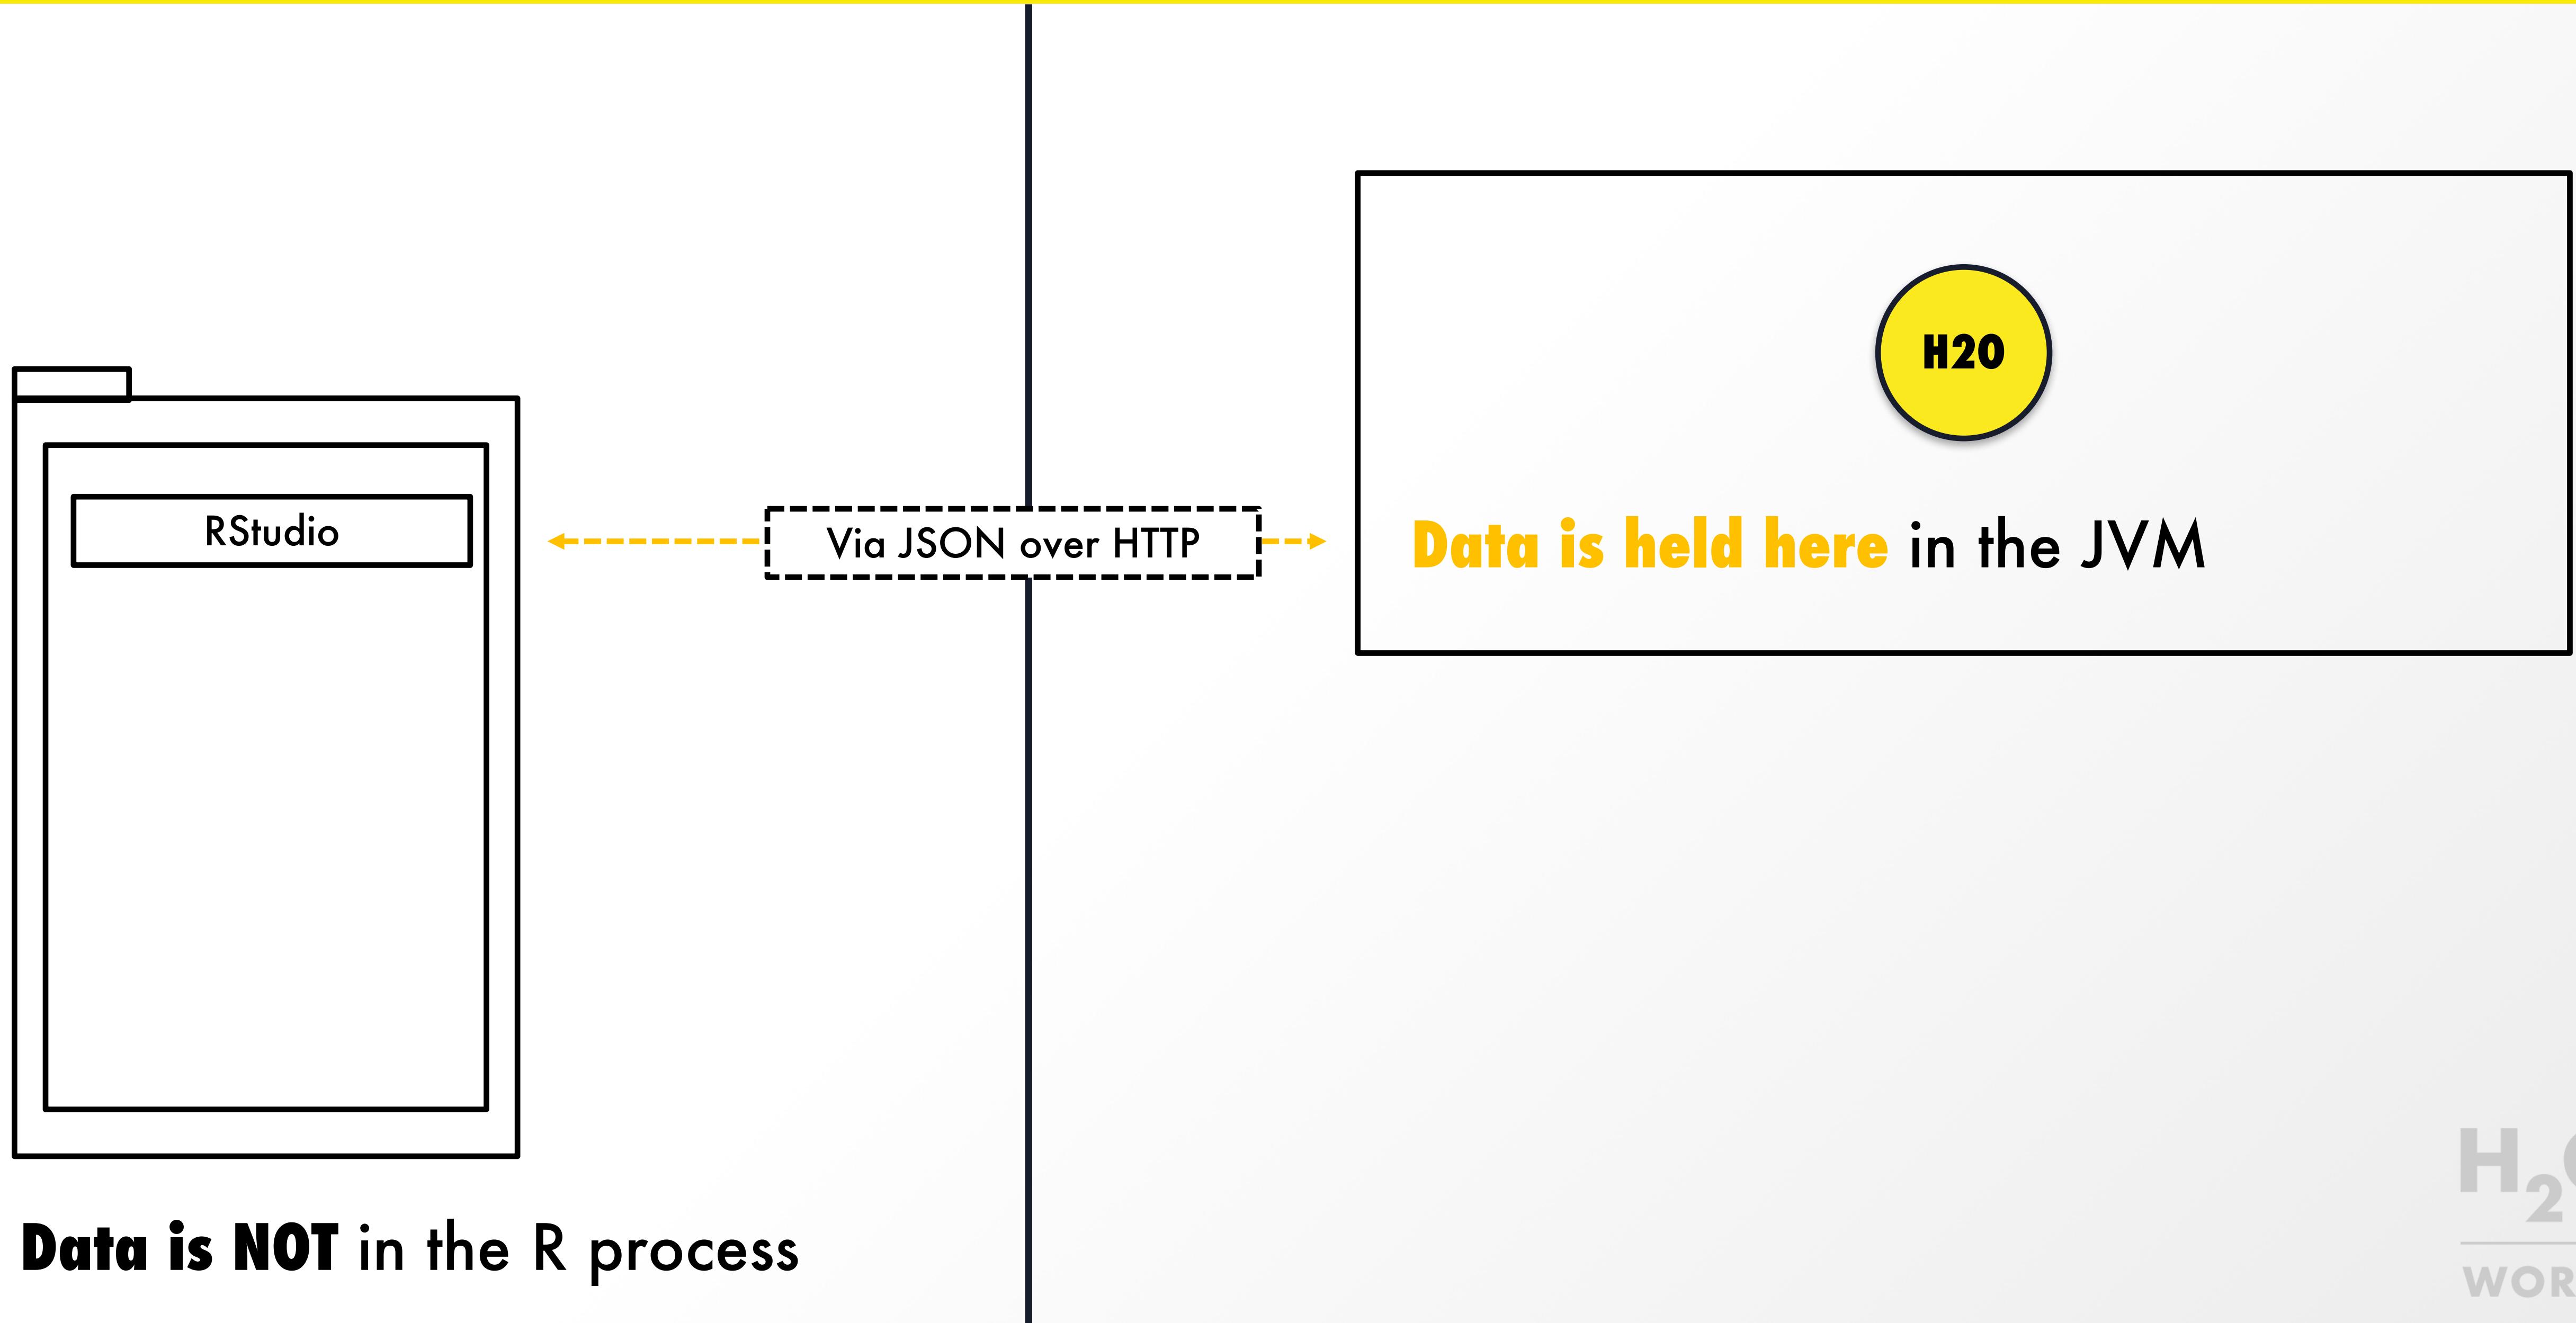
# H2O R API



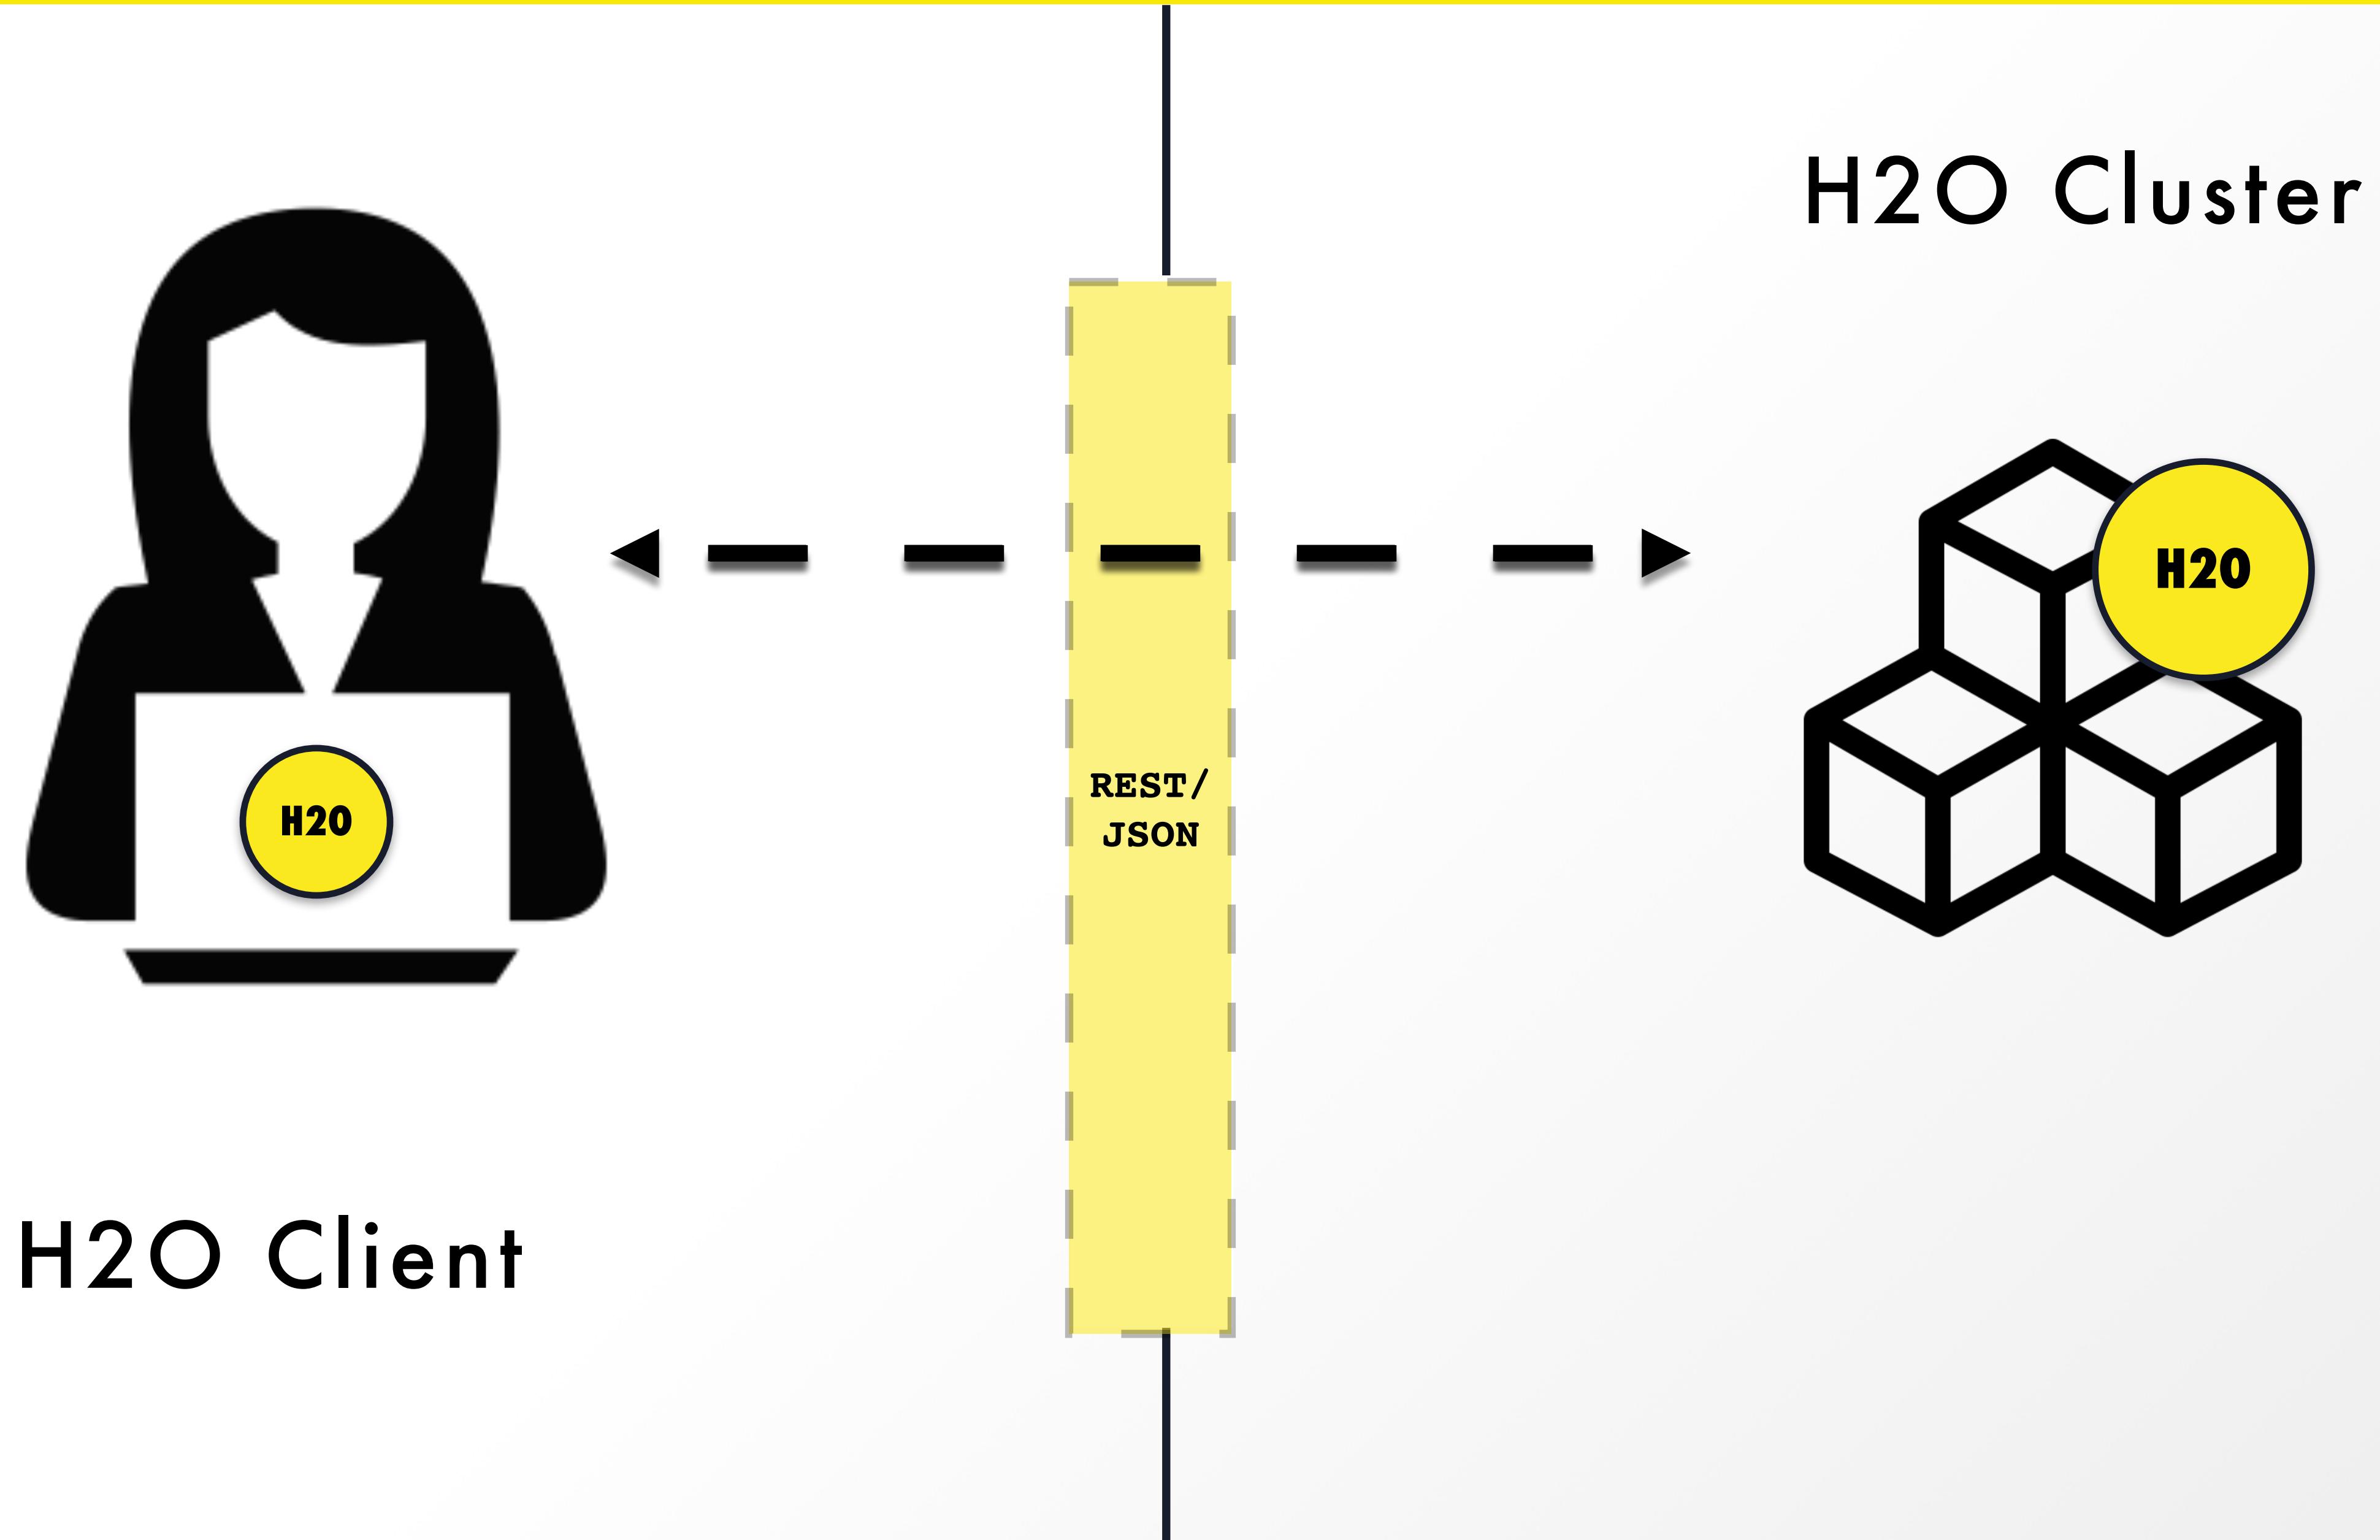
# H2O R API



# H2O R API



# Client & Cluster Communication



# Communication Layers: Interface

Standard R  
Interface



the Client

RStudio Interface

```
Console | Terminal ~ /Desktop/rencontres-R-2018/h2o-deeplearning/ ~

H2O is not running yet, starting it now...

Note: In case of errors look at the following log files:
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
  H2O cluster uptime:      3 seconds 970 milliseconds
  H2O cluster timezone:    America/Los_Angeles
  H2O data parsing timezone: UTC
  H2O cluster version:     3.20.0.2
  H2O cluster version age: 12 days
  H2O cluster name:        H2O_started_from_R_navdeepgill_kdm352
  H2O cluster total nodes: 1
  H2O cluster total memory: 3.56 GB
  H2O cluster total cores: 8
  H2O cluster allowed cores: 8
  H2O cluster healthy:    TRUE
  H2O Connection ip:       localhost
  H2O Connection port:     54321
  H2O Connection proxy:   NA
  H2O Internal Security: FALSE
  H2O API Extensions:    XGBoost, Algos, AutoML, Core V3, Core V4
  R Version:              R version 3.4.0 (2017-04-21)
```

# Communication Layers: Code Script

## Rstudio using H2O Package

RStudio



the Client

```
Console Terminal ~~/Desktop/rencontres-R-2018/h2o-deeplearning/ ~

H2O is not running yet, starting it now...

Note: In case of errors look at the following log files:
/var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
/var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
H2O cluster uptime:      3 seconds 970 milliseconds
H2O cluster timezone:    America/Los_Angeles
H2O data parsing timezone: UTC
H2O cluster version:     3.20.0.2
H2O cluster version age: 12 days
H2O cluster name:        H2O_started_from_R_navdeepgill_kdm352
H2O cluster total nodes: 1
H2O cluster total memory: 3.56 GB
H2O cluster total cores: 8
H2O cluster allowed cores: 8
H2O cluster healthy:     TRUE
H2O Connection ip:       localhost
H2O Connection port:     54321
H2O Connection proxy:    NA
H2O Internal Security:  FALSE
H2O API Extensions:    XGBoost, Algos, AutoML, Core V3, Core V4
R Version:               R version 3.4.0 (2017-04-21)

> |
```

# Communication Layers: A Command

## Importing Big Data with R Code

R Commands



the Client

```
Console Terminal ~~/Desktop/rencontres-R-2018/h2o-deeplearning/ ~

H2O is not running yet, starting it now...

Note: In case of errors look at the following log files:
/var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
/var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

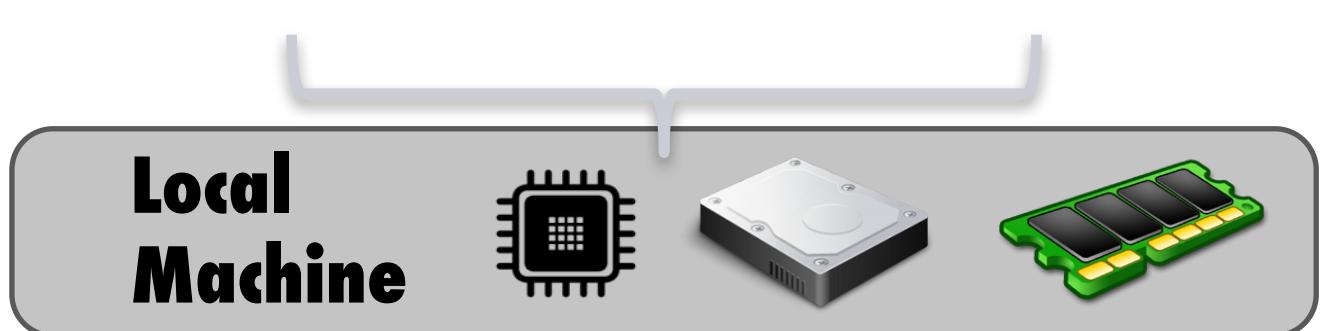
java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
  H2O cluster uptime:      3 seconds 970 milliseconds
  H2O cluster timezone:    America/Los_Angeles
  H2O data parsing timezone: UTC
  H2O cluster version:    3.20.0.2
  H2O cluster version age: 12 days
  H2O cluster name:       H2O_started_from_R_navdeepgill_kdm352
  H2O cluster total nodes: 1
  H2O cluster total memory: 3.56 GB
  H2O cluster total cores: 8
  H2O cluster allowed cores: 8
  H2O cluster healthy:     TRUE
  H2O Connection ip:       localhost
  H2O Connection port:     54321
  H2O Connection proxy:   NA
  H2O Internal Security:  FALSE
  H2O API Extensions:    XGBoost, Algos, AutoML, Core V3, Core V4
  R Version:              R version 3.4.0 (2017-04-21)

> |
```

**h2o.importFile(...)**



# Communication Layers: A Command

## Importing Big Data with R Code

R Commands



the Client

```
Console Terminal ~ /Desktop/rencontres-R-2018/h2o-deeplearning/ ~
H2O is not running yet, starting it now...
Note: In case of errors look at the following log files:
      /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
      /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

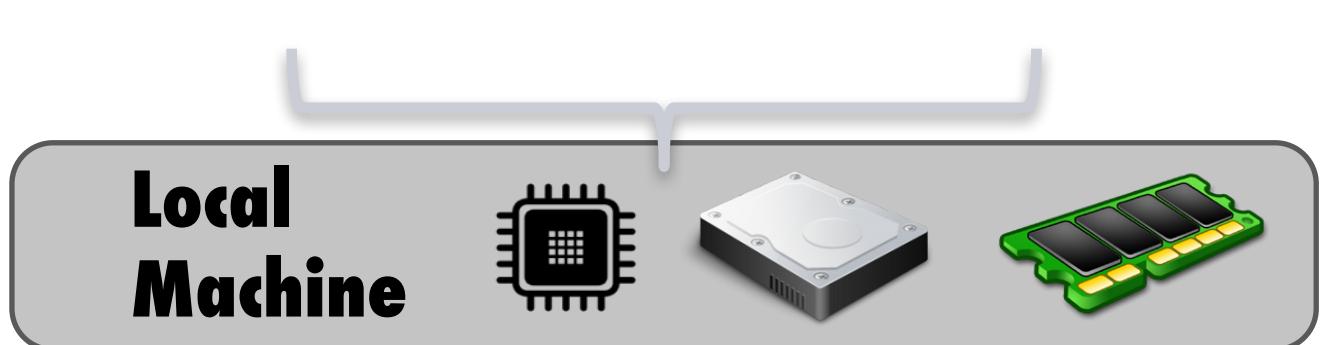
Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
  H2O cluster uptime:      3 seconds 970 milliseconds
  H2O cluster timezone:    America/Los_Angeles
  H2O data parsing timezone: UTC
  H2O cluster version:    3.20.0.2
  H2O cluster version age: 12 days
  H2O cluster name:       H2O_started_from_R_navdeepgill_kdm352
  H2O cluster total nodes: 1
  H2O cluster total memory: 3.56 GB
  H2O cluster total cores: 8
  H2O cluster allowed cores: 8
  H2O cluster healthy:    TRUE
  H2O Connection ip:      localhost
  H2O Connection port:    54321
  H2O Connection proxy:   NA
  H2O Internal Security: FALSE
  H2O API Extensions:   XGBoost, Algos, AutoML, Core V3, Core V4
  R Version:              R version 3.4.0 (2017-04-21)

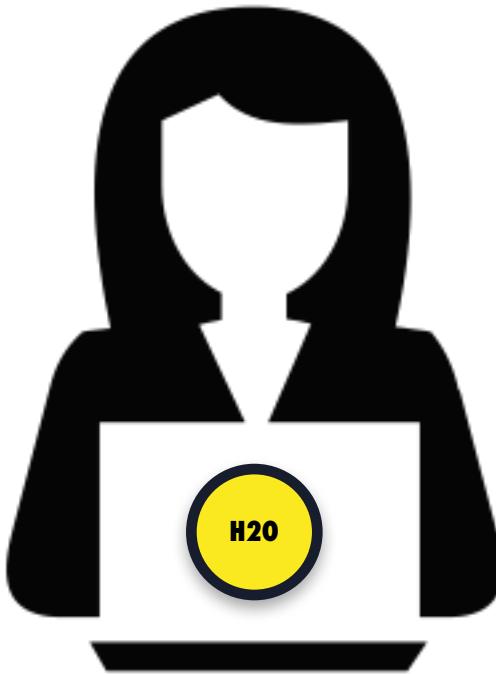
> |
```

Path to Your Dataset

`h2o.importFile(...)`



# Fourth: Communicate



`h2o.importFile(...)`  
requests file import

```
Console Terminal ~Desktop/rencontres-R-2018/h2o-deeplearning/ ~

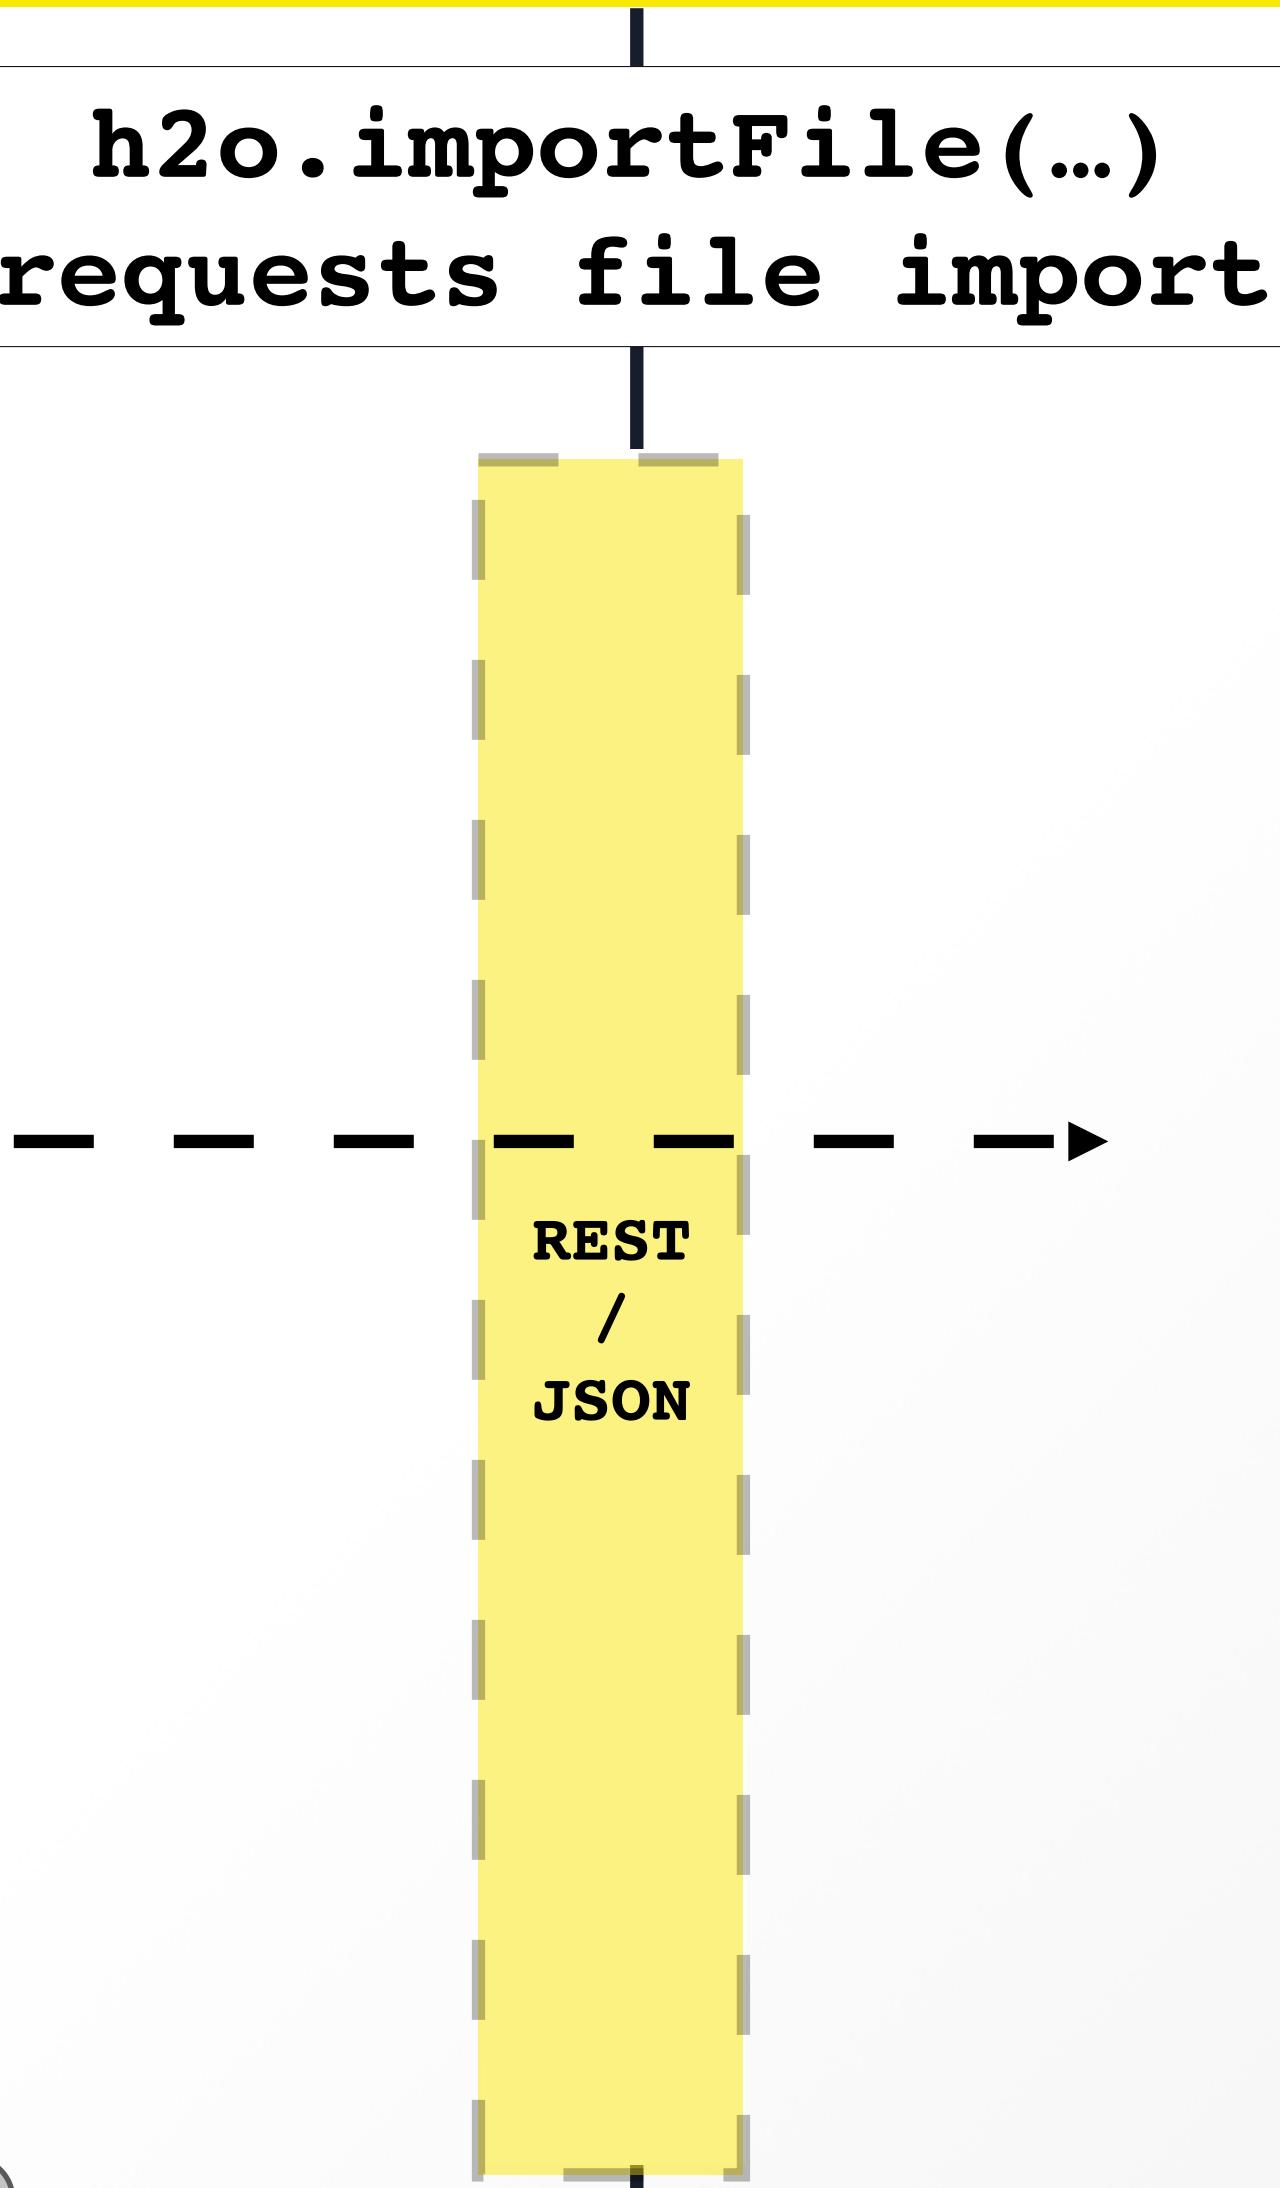
H2O is not running yet, starting it now...

Note: In case of errors look at the following log files:
  /var/folders/55/rj4cny_s29q4vn1jt_x08sm0000gn/T/RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
  /var/folders/55/rj4cny_s29q4vn1jt_x08sm0000gn/T/RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
H2O cluster uptime: 3 seconds 970 milliseconds
H2O cluster timezone: America/Los_Angeles
H2O data parsing timezone: UTC
H2O cluster version: 3.28.0.2
H2O cluster version age: 12 days
H2O cluster name: H2O_started_from_R_navdeepgill_kdm352
H2O cluster total nodes: 1
H2O cluster total memory: 3.56 GB
H2O cluster total cores: 8
H2O cluster allowed cores: 8
H2O cluster healthy: TRUE
H2O Connection ip: localhost
H2O Connection port: 54321
H2O Connection proxy: NA
H2O Internal Security: FALSE
H2O API Extensions: XGBoost, Algos, AutoML, Core V3, Core V4
R Version: R version 3.4.0 (2017-04-21)
```



H2O Cluster



# Fifth: Cluster Does Heavy Lifting



```
Console Terminal ~ /Desktop/rencontres-R-2018/h2o-deeplearning/ ~

H2O is not running yet, starting it now...

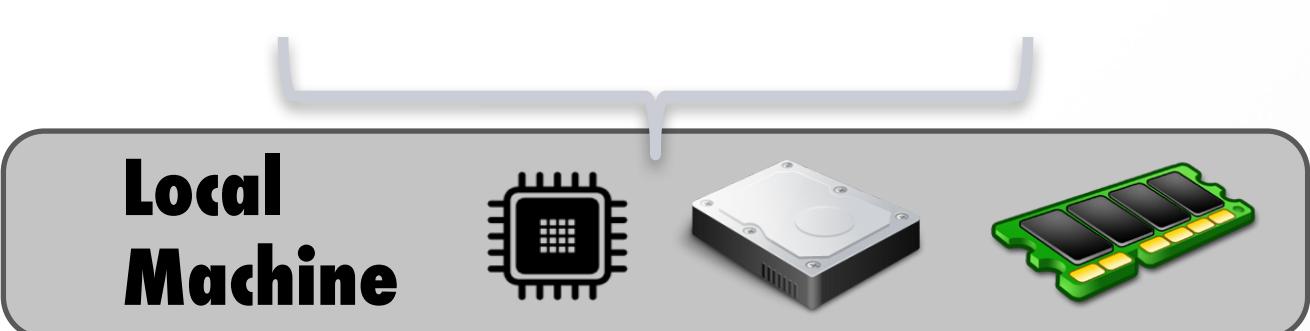
Note: In case of errors look at the following log files:
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZlxR/h2o_navdeepgill_started_from_r.out
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZlxR/h2o_navdeepgill_started_from_r.err

java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
H2O cluster uptime: 3 seconds 970 milliseconds
H2O cluster timezone: America/Los_Angeles
H2O data parsing timezone: UTC
H2O cluster version: 3.20.0.2
H2O cluster version age: 12 days
H2O cluster name: H2O_started_from_R_navdeepgill_kdm352
H2O cluster total nodes: 1
H2O cluster total memory: 3.56 GB
H2O cluster total cores: 8
H2O cluster allowed cores: 8
H2O cluster healthy: TRUE
H2O Connection ip: localhost
H2O Connection port: 54321
H2O Connection proxy: NA
H2O Internal Security: FALSE
H2O API Extensions: XGBoost, Algos, AutoML, Core V3, Core V4
R Version: R version 3.4.0 (2017-04-21)

> |
```



H2O Cluster

REST/  
JSON

# Fifth: Cluster Does Heavy Lifting



```
Console Terminal ~/Desktop/rencontres-R-2018/h2o-deeplearning/ ~

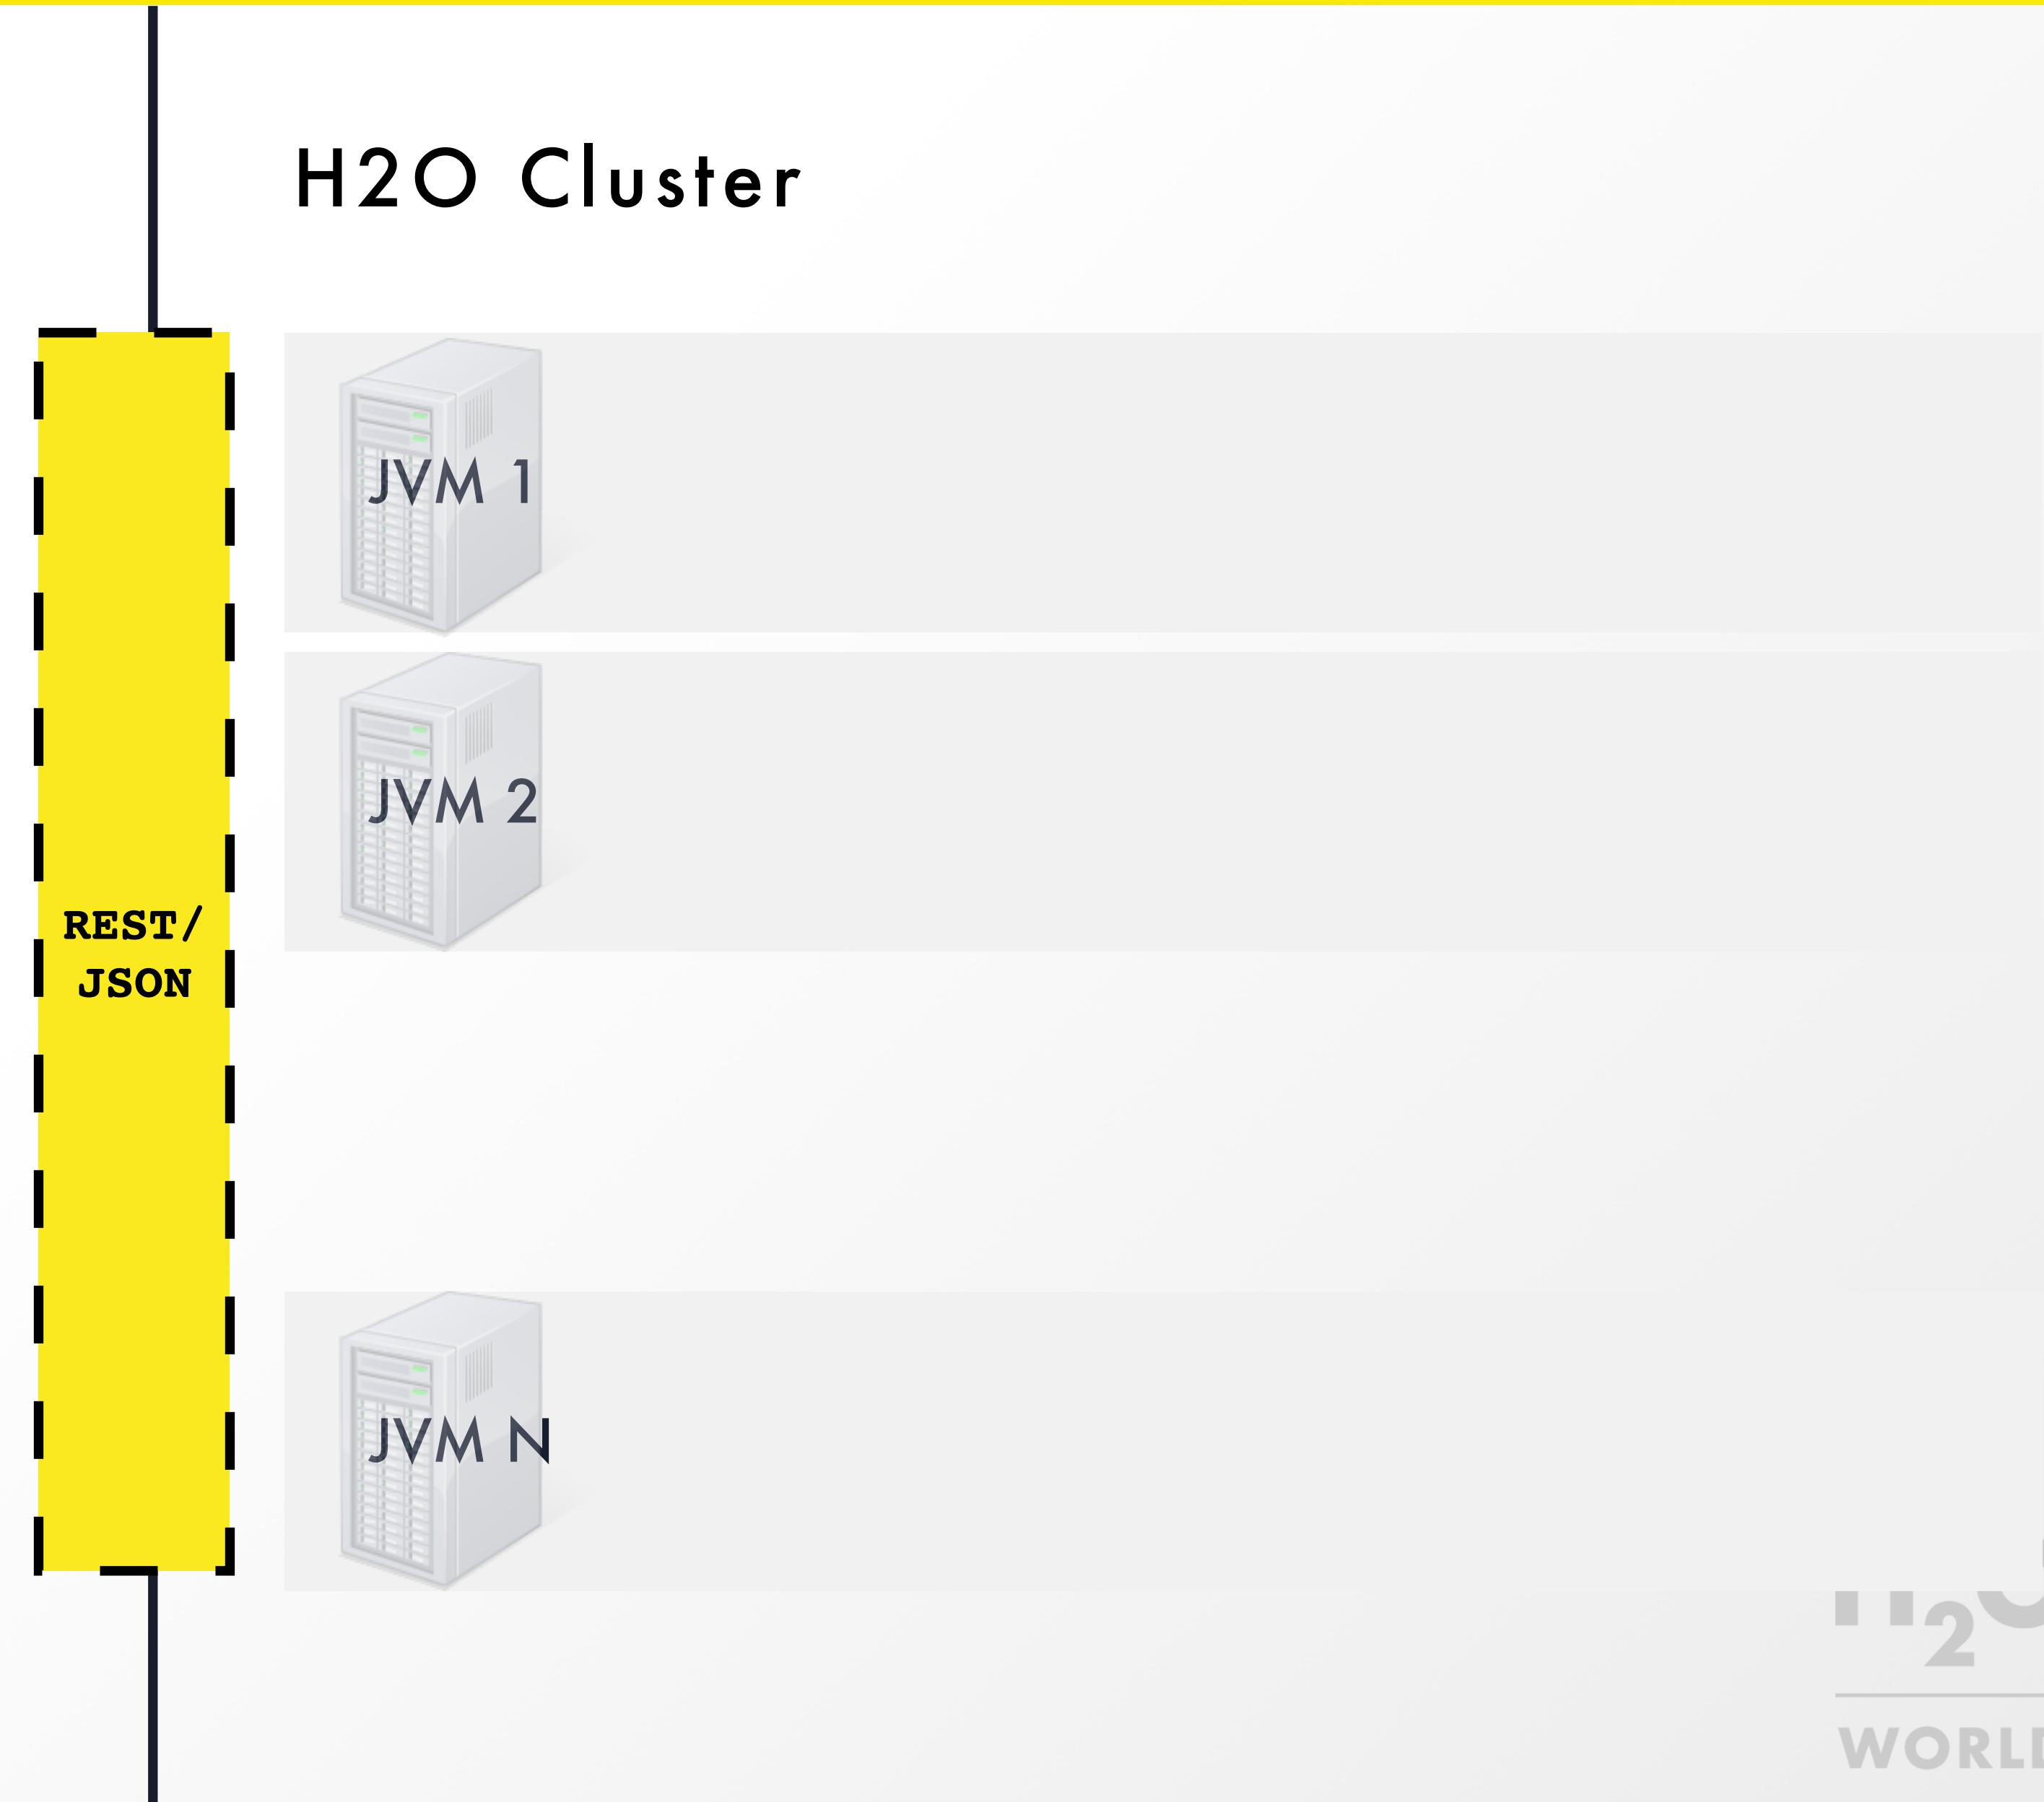
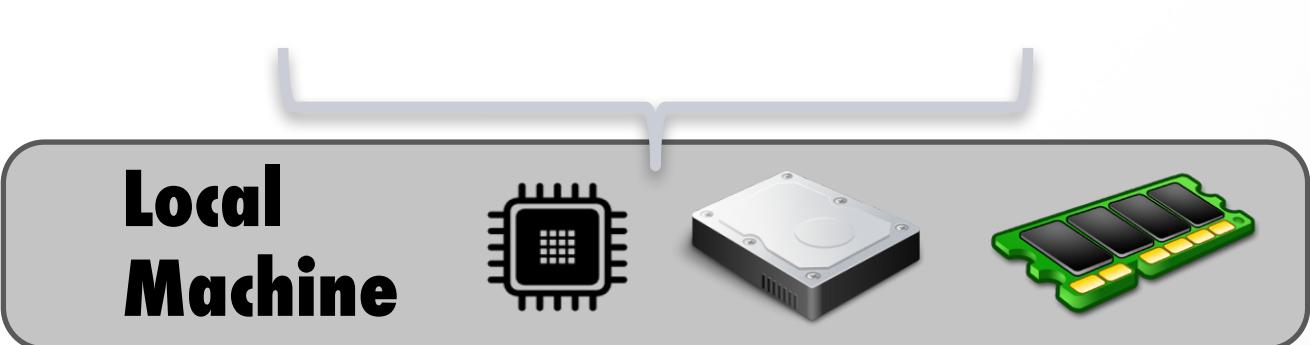
H2O is not running yet, starting it now...

Note: In case of errors look at the following log files:
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

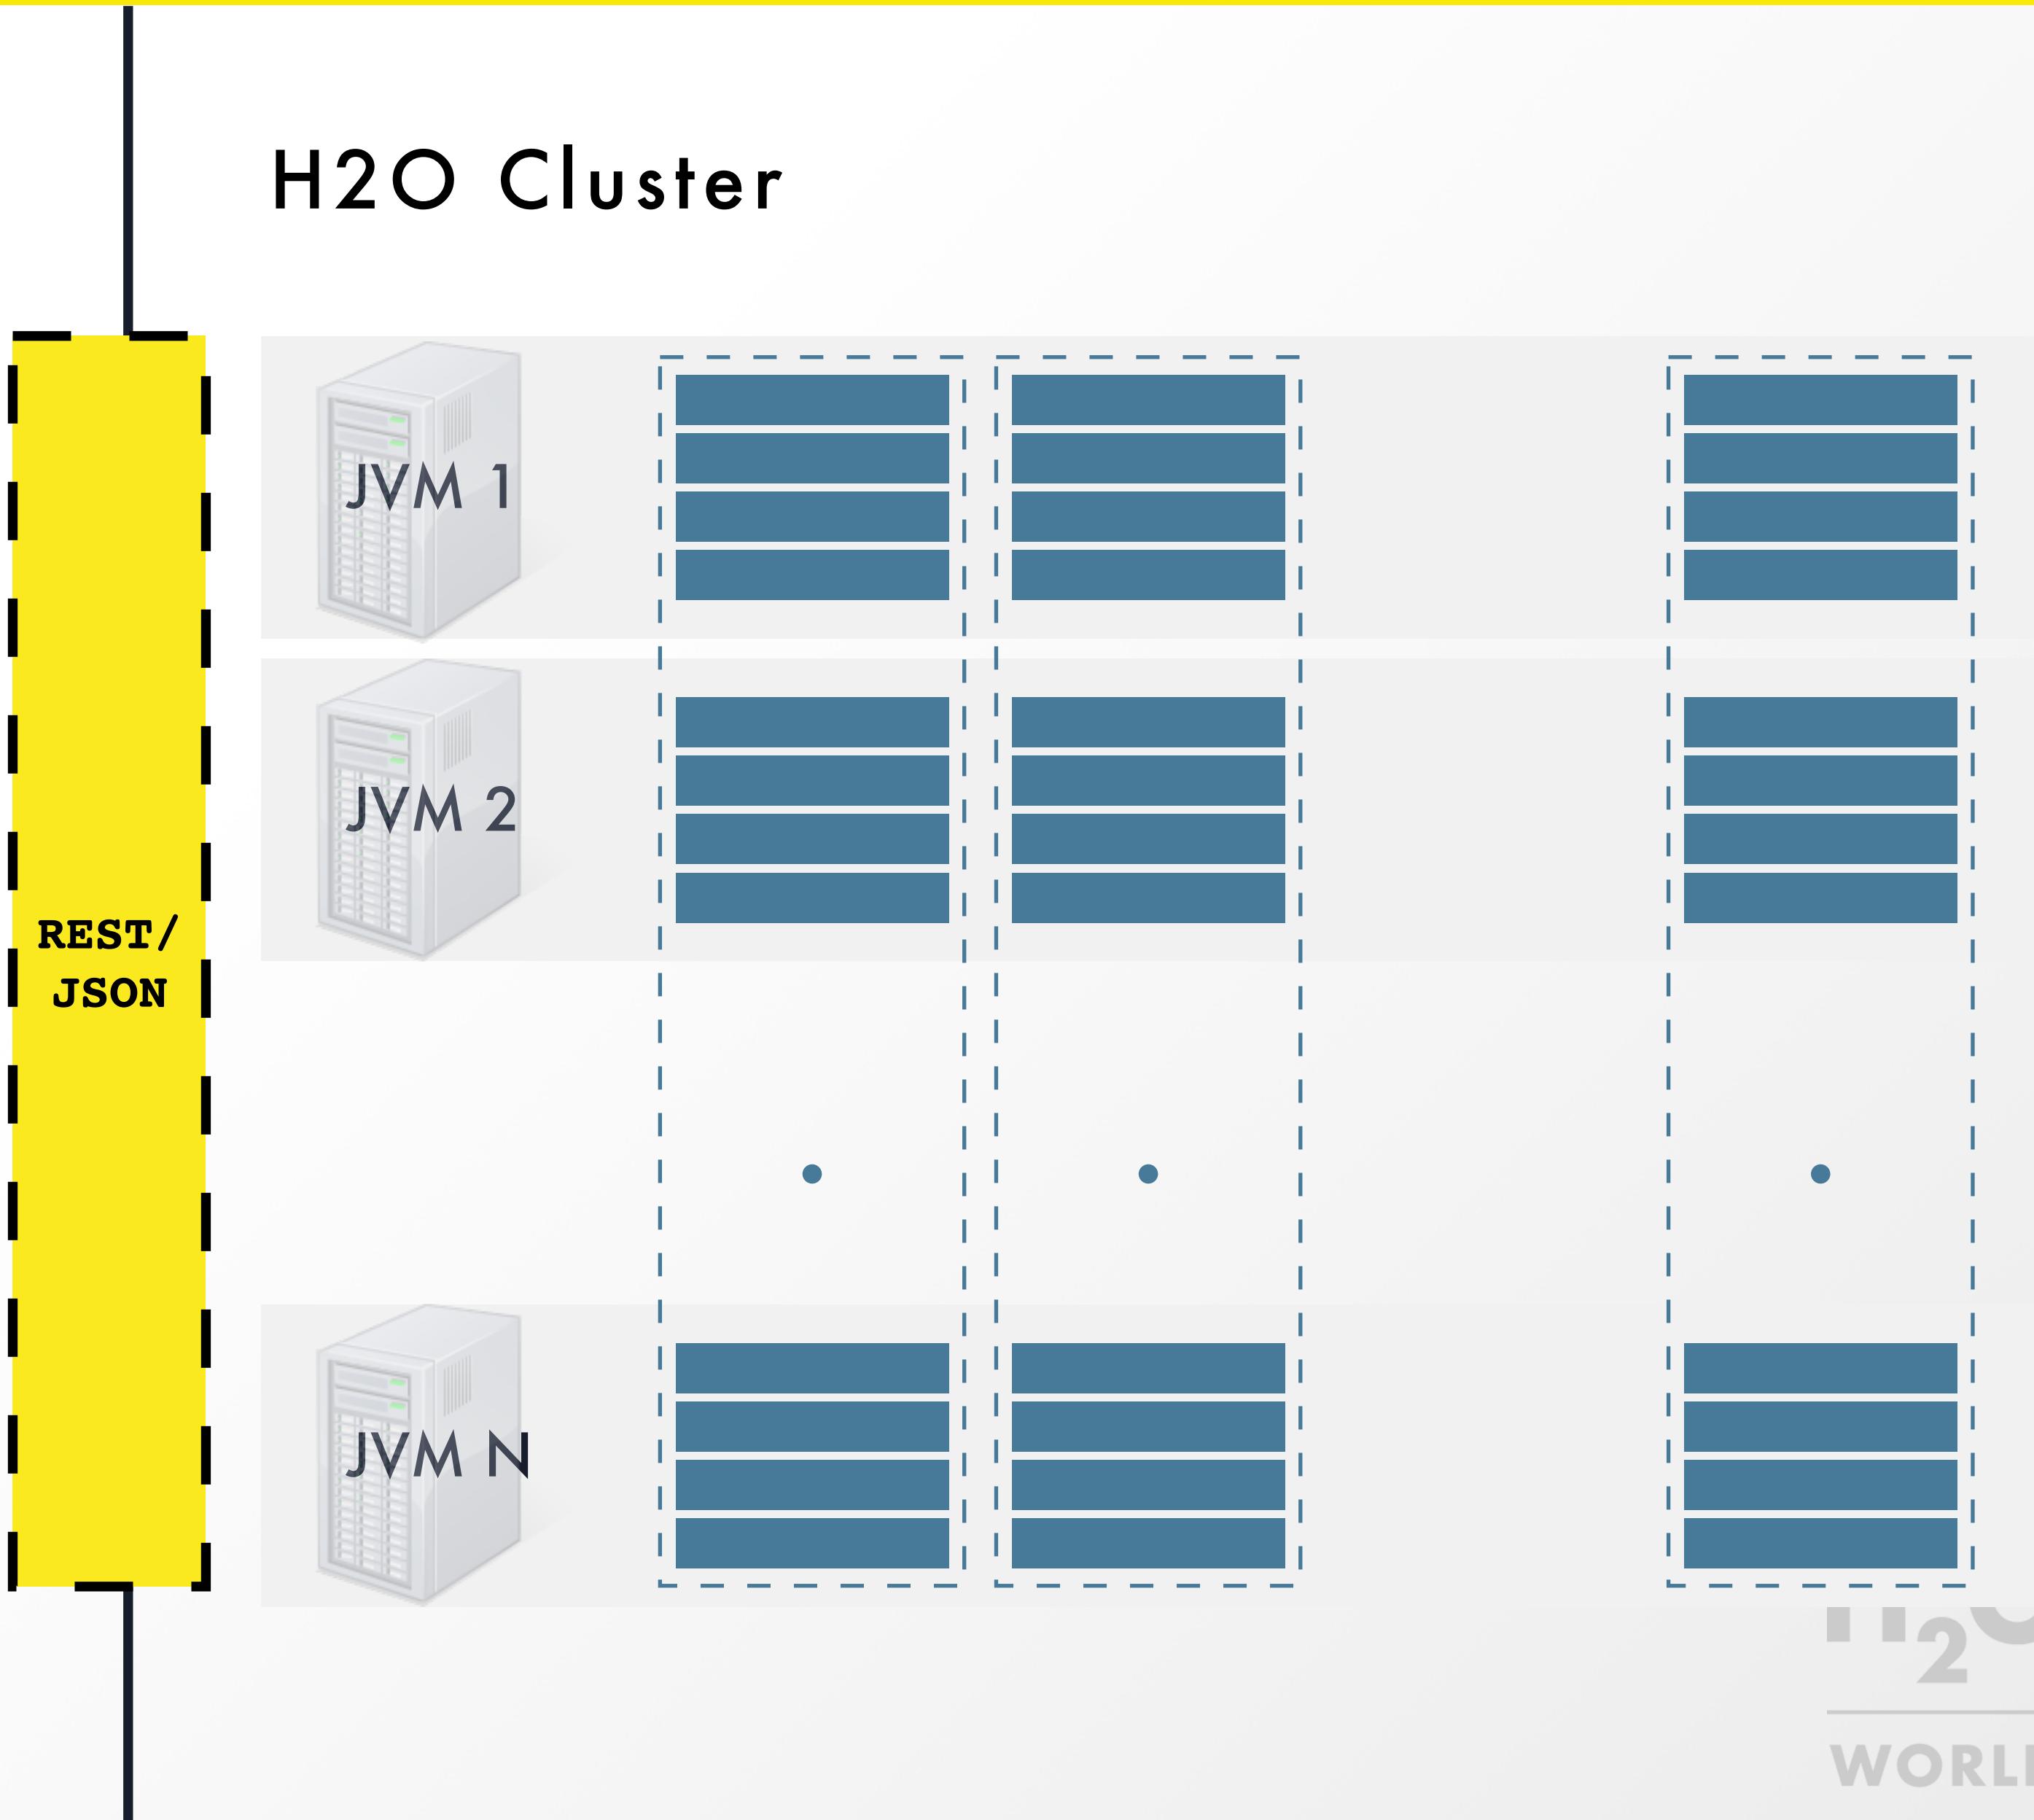
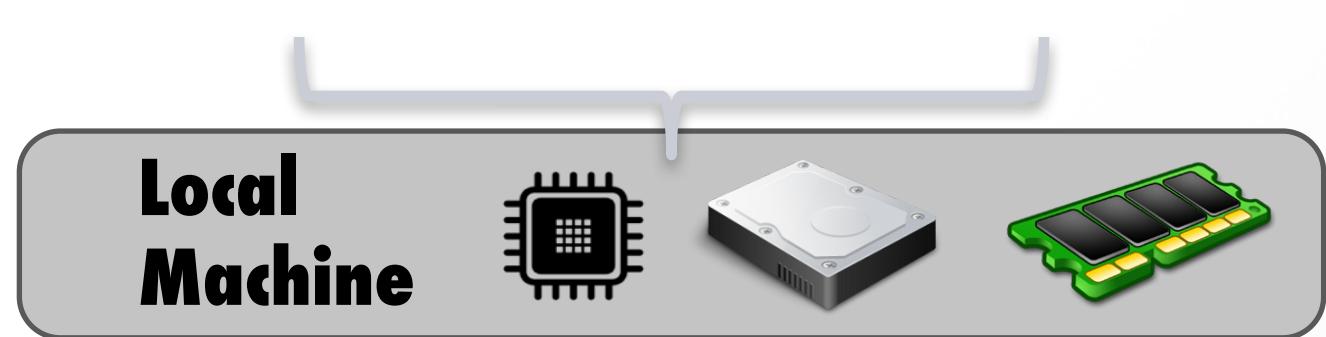
R is connected to the H2O cluster:
  H2O cluster uptime:      3 seconds 970 milliseconds
  H2O cluster timezone:    America/Los_Angeles
  H2O data parsing timezone: UTC
  H2O cluster version:     3.20.0.2
  H2O cluster version age: 12 days
  H2O cluster name:        H2O_started_from_R_navdeepgill_kdm352
  H2O cluster total nodes: 1
  H2O cluster total memory: 3.56 GB
  H2O cluster total cores: 8
  H2O cluster allowed cores: 8
  H2O cluster healthy:     TRUE
  H2O Connection ip:       localhost
  H2O Connection port:     54321
  H2O Connection proxy:    NA
  H2O Internal Security:   FALSE
  H2O API Extensions:     XGBoost, Algos, AutoML, Core V3, Core V4
  R Version:               R version 3.4.0 (2017-04-21)
```



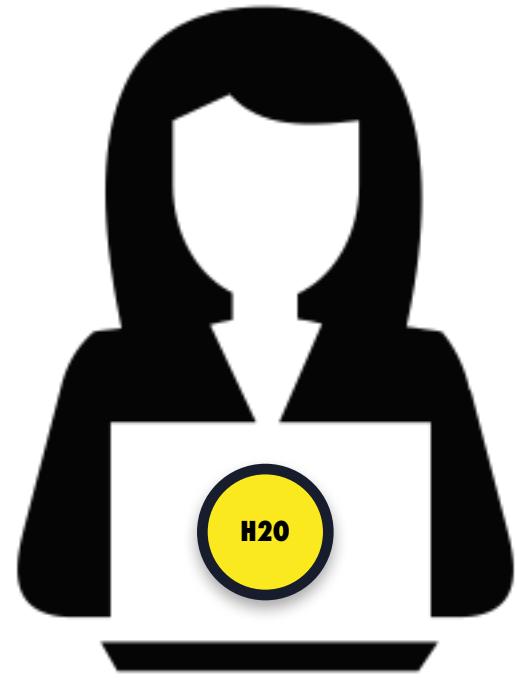
# Fifth: Cluster Does Heavy Lifting



```
Console Terminal ~/Desktop/rencontres-R-2018/h2o-deeplearning/ ~  
H2O is not running yet, starting it now...  
  
Note: In case of errors look at the following log files:  
/var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZlxR/h2o_navdeepgill_started_from_r.out  
/var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZlxR/h2o_navdeepgill_started_from_r.err  
  
java version "1.8.0_101"  
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)  
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)  
  
Starting H2O JVM and connecting: . Connection successful!  
  
R is connected to the H2O cluster:  
H2O cluster uptime: 3 seconds 970 milliseconds  
H2O cluster timezone: America/Los_Angeles  
H2O data parsing timezone: UTC  
H2O cluster version: 3.20.0.2  
H2O cluster version age: 12 days  
H2O cluster name: H2O_started_from_R_navdeepgill_kdm352  
H2O cluster total nodes: 1  
H2O cluster total memory: 3.56 GB  
H2O cluster total cores: 8  
H2O cluster allowed cores: 8  
H2O cluster healthy: TRUE  
H2O Connection ip: localhost  
H2O Connection port: 54321  
H2O Connection proxy: NA  
H2O Internal Security: FALSE  
H2O API Extensions: XGBoost, Algos, AutoML, Core V3, Core V4  
R Version: R version 3.4.0 (2017-04-21)  
> |
```



# Fifth: Cluster Does Heavy Lifting



```
Console Terminal ~ /Desktop/rencontres-R-2018/h2o-deeplearning/ ↵

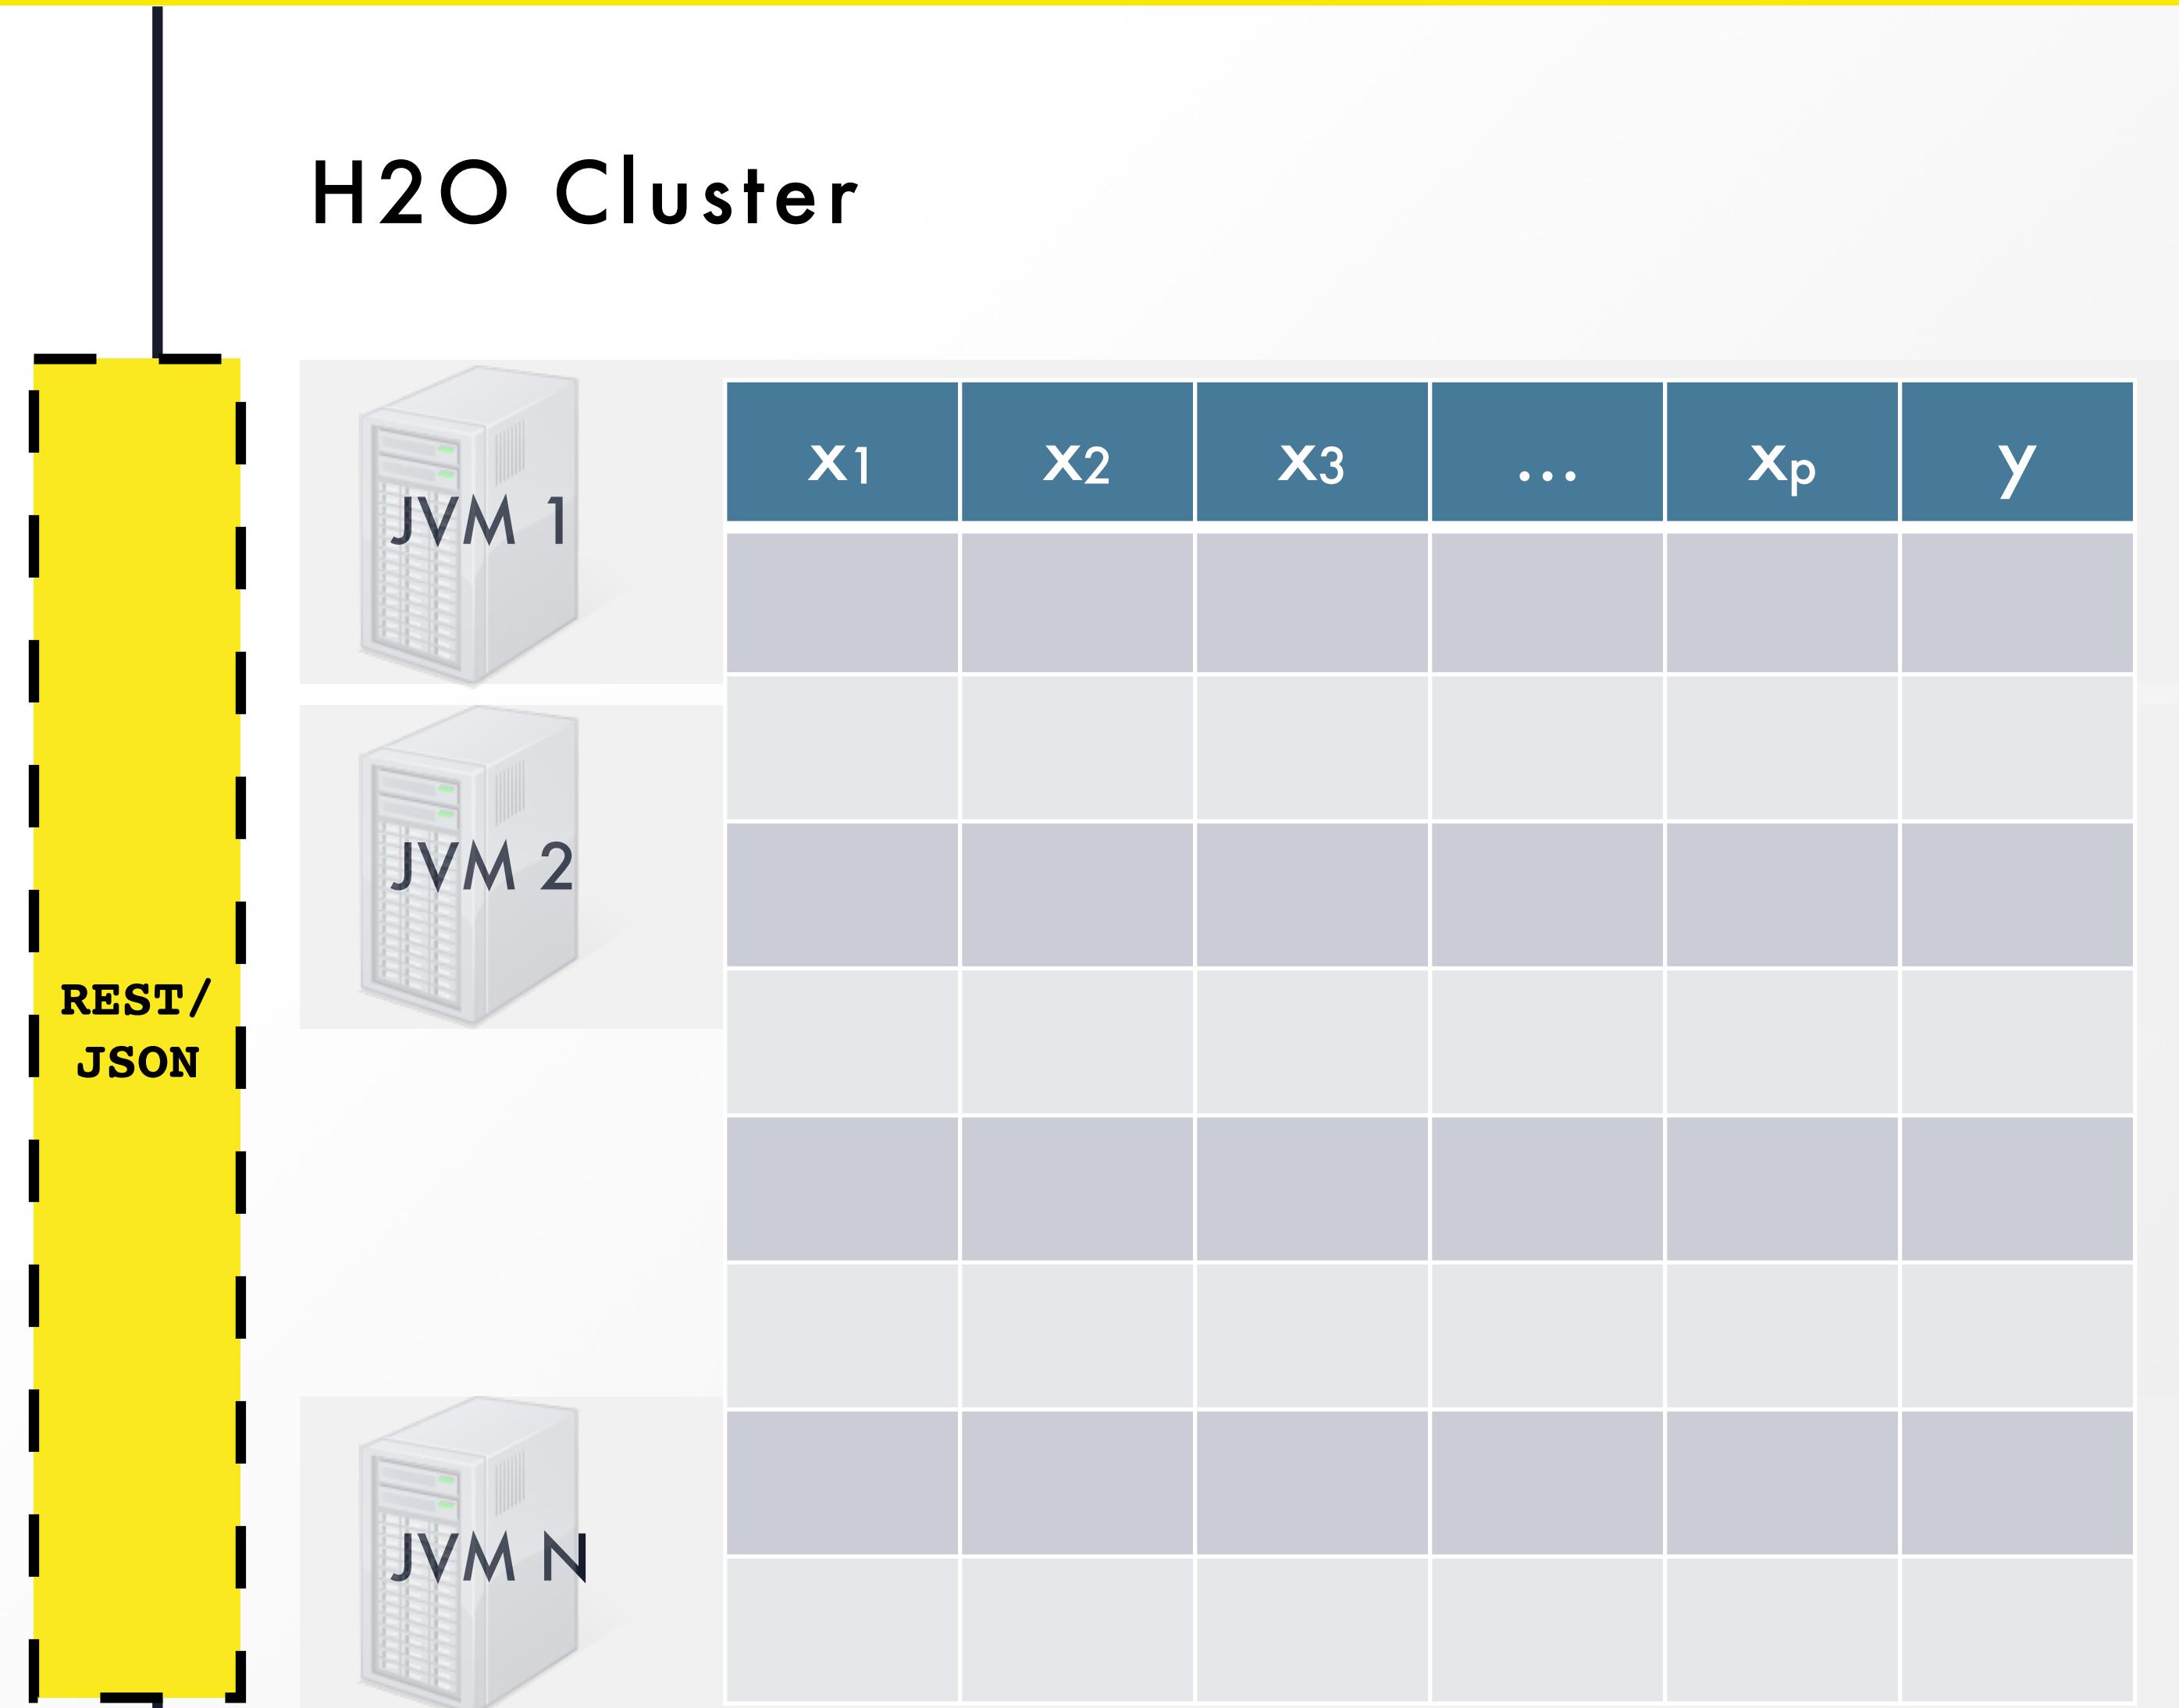
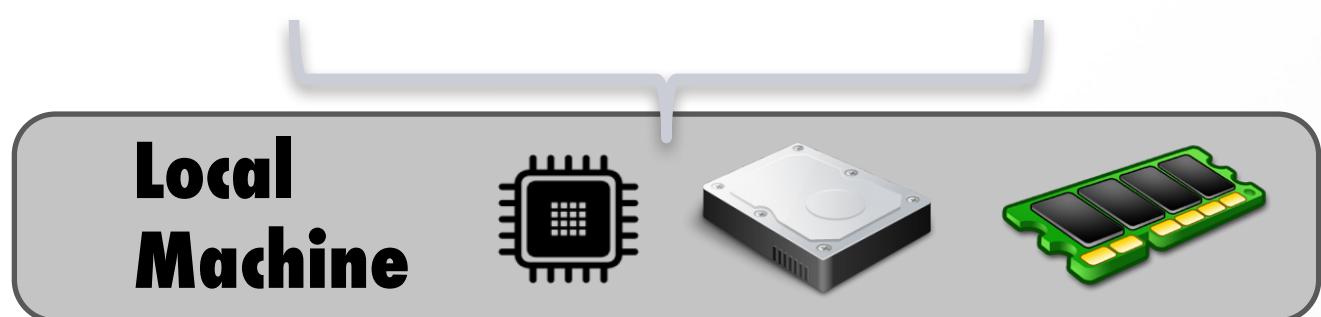
H2O is not running yet, starting it now...

Note: In case of errors look at the following log files:
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.out
  /var/folders/55/rj4cny_s29q4vn1wjt_x08sm0000gn/T//RtmpH6ZkxR/h2o_navdeepgill_started_from_r.err

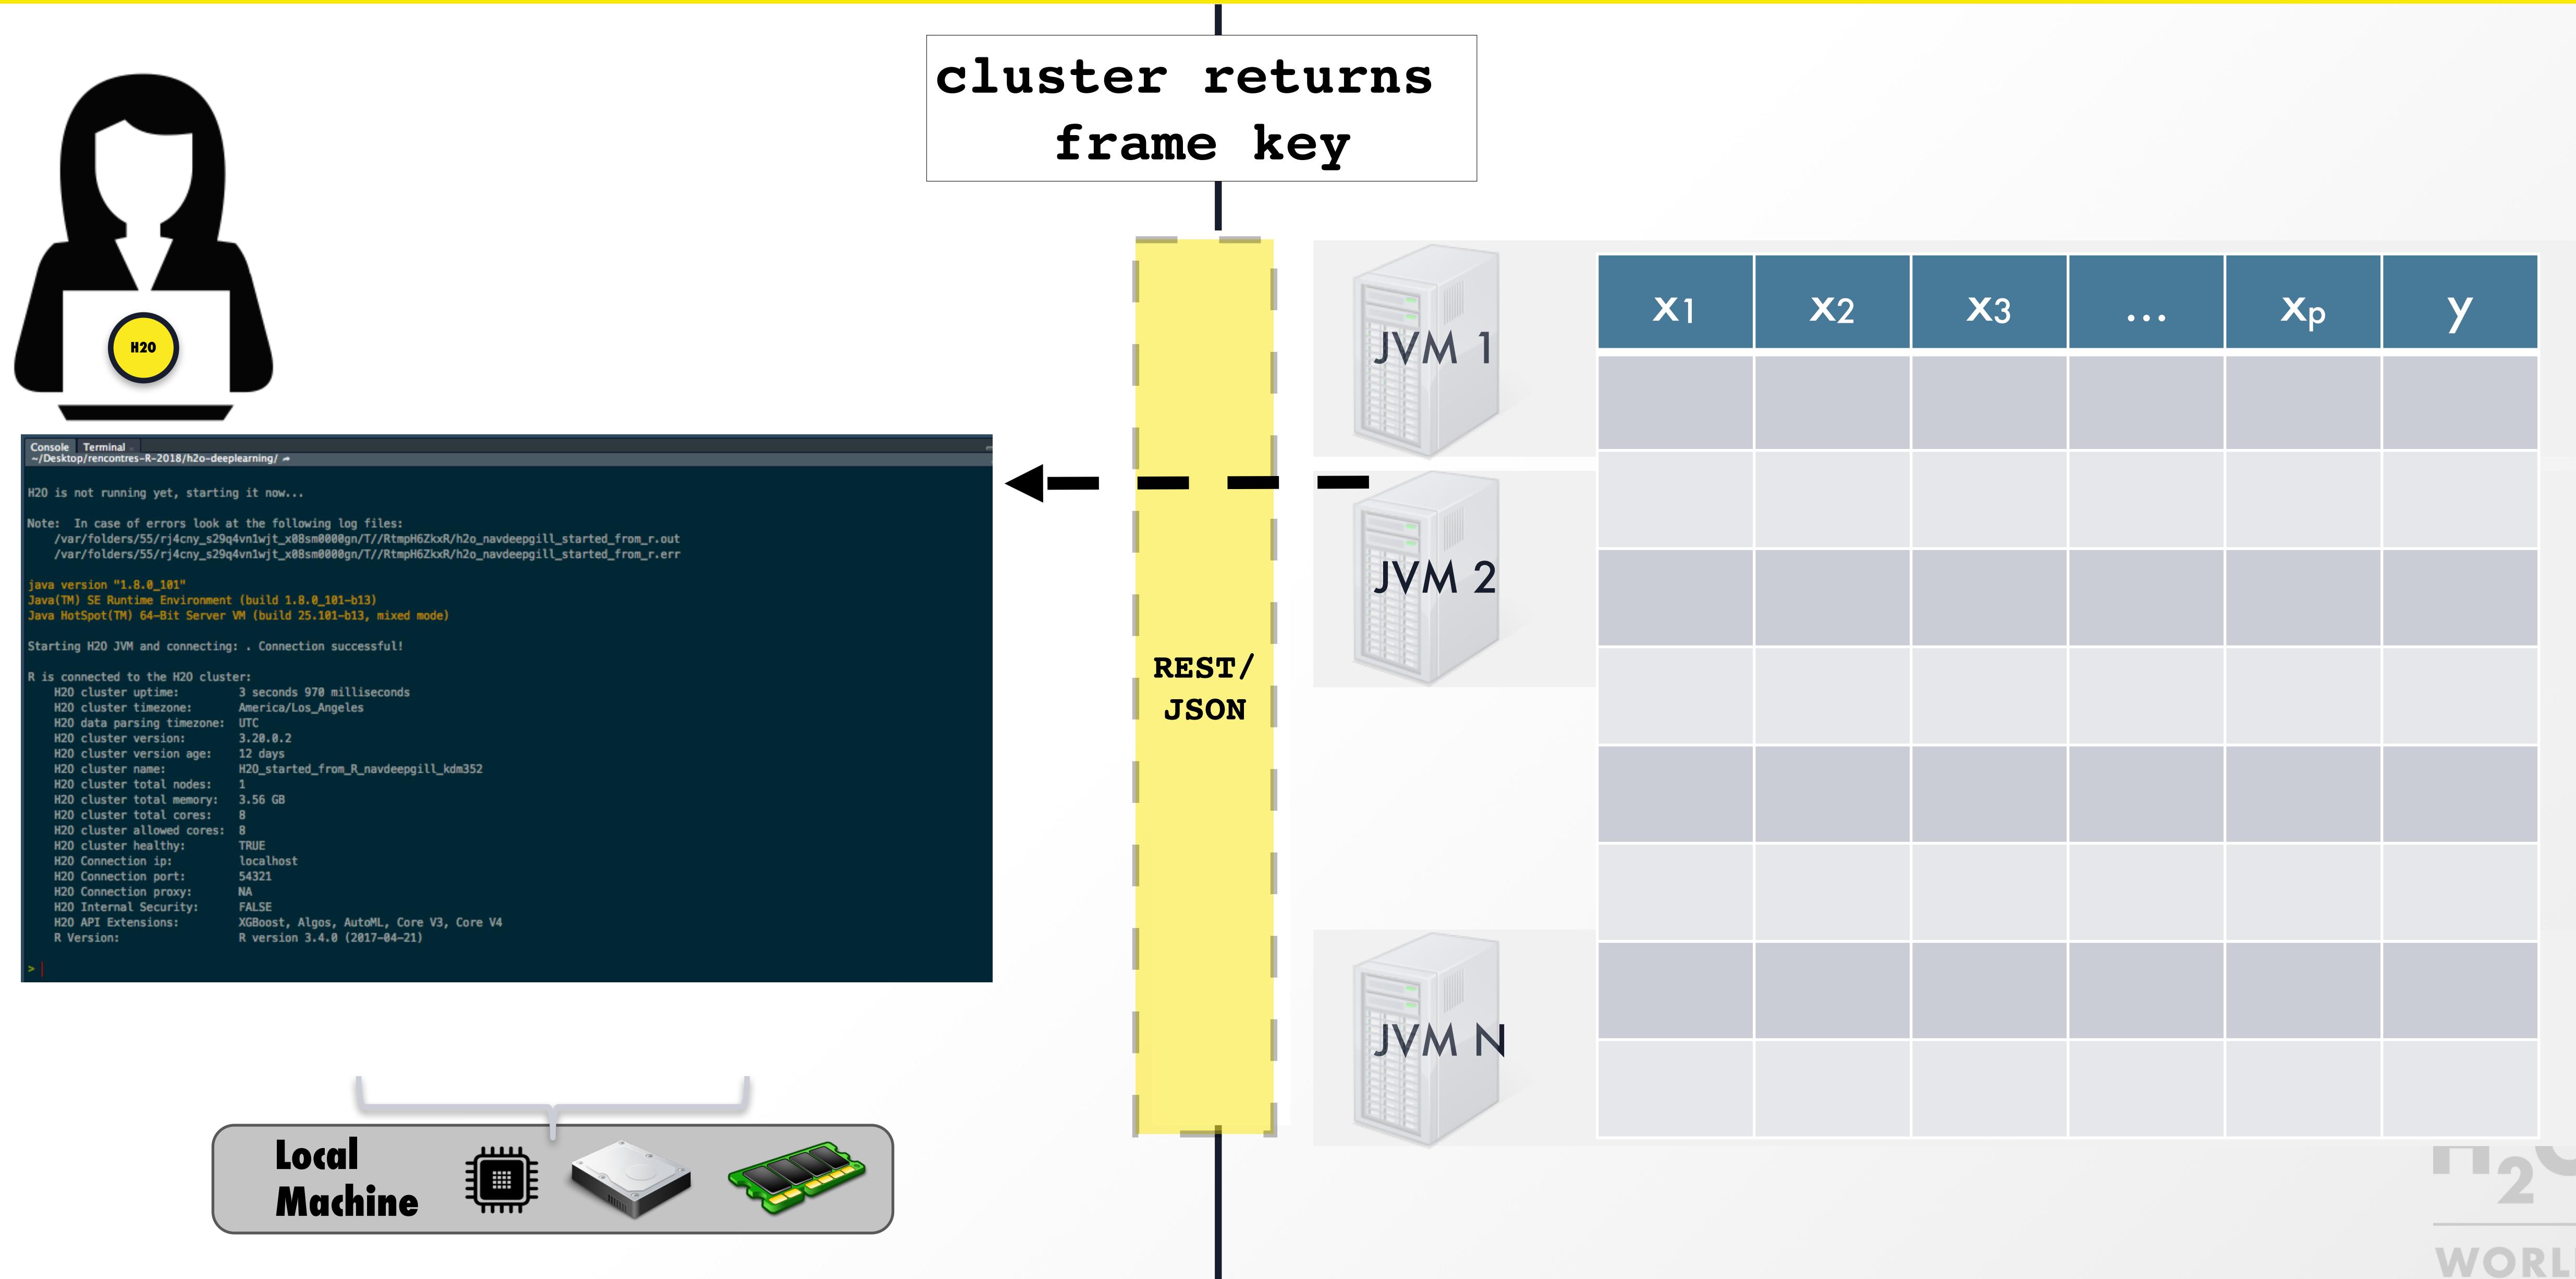
java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)

Starting H2O JVM and connecting: . Connection successful!

R is connected to the H2O cluster:
  H2O cluster uptime:      3 seconds 970 milliseconds
  H2O cluster timezone:    America/Los_Angeles
  H2O data parsing timezone: UTC
  H2O cluster version:    3.20.0.2
  H2O cluster version age: 12 days
  H2O cluster name:       H2O_started_from_R_navdeepgill_kdm352
  H2O cluster total nodes: 1
  H2O cluster total memory: 3.56 GB
  H2O cluster total cores: 8
  H2O cluster allowed cores: 8
  H2O cluster healthy:     TRUE
  H2O Connection ip:       localhost
  H2O Connection port:     54321
  H2O Connection proxy:    NA
  H2O Internal Security:  FALSE
  H2O API Extensions:    XGBoost, Algos, AutoML, Core V3, Core V4
  R Version:              R version 3.4.0 (2017-04-21)
```

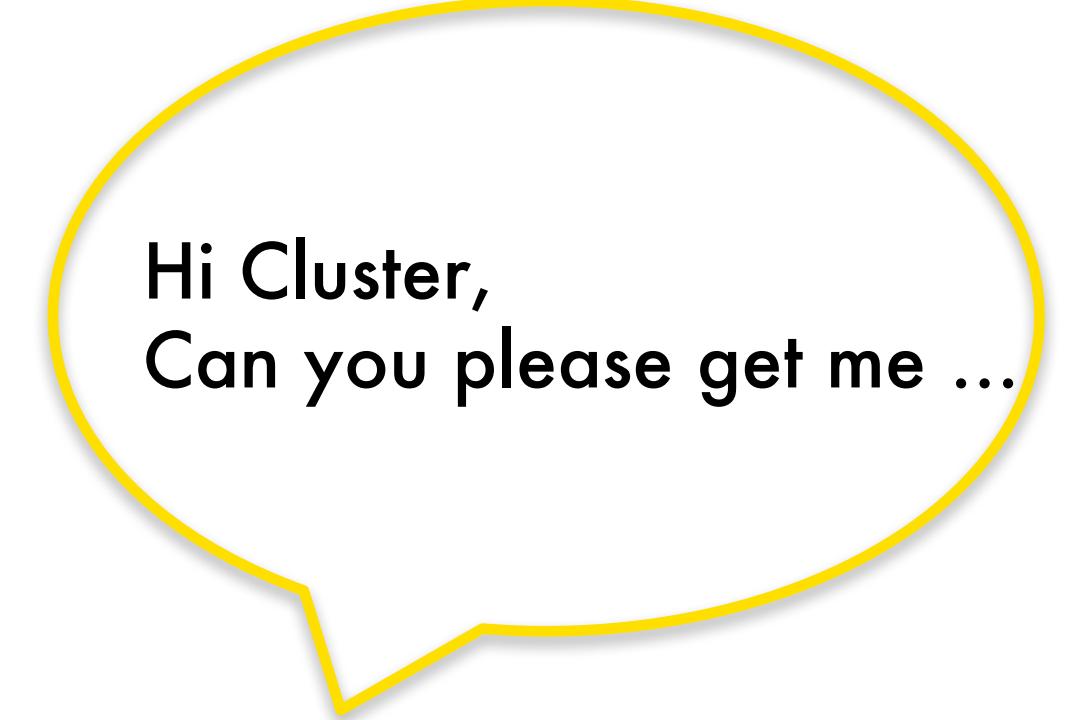


# H2O Client



# What does the Client Do?

Clients Only Tells the Cluster What to Do

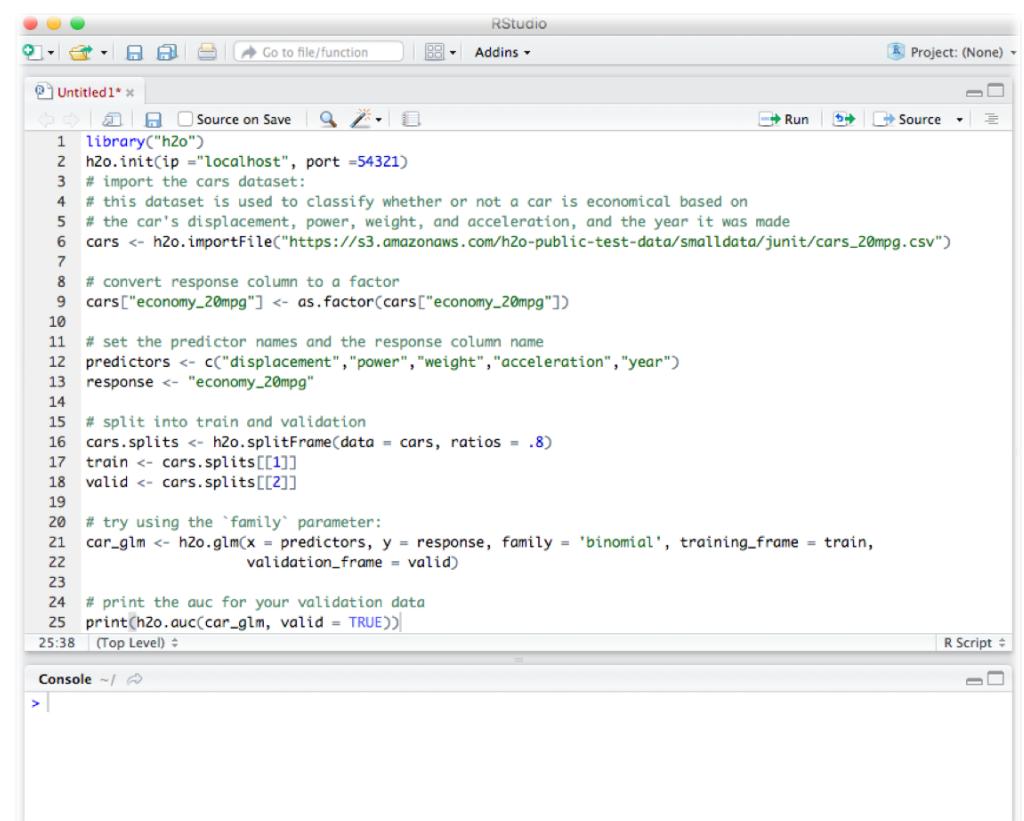


```
library("h2o")
h2o.init(ip = "localhost", port = 54321)
# import the cars dataset:
# this dataset is used to classify whether or not a car is economical based on
# the car's displacement, power, weight, and acceleration, and the year it was made
cars <- h2o.importFile("https://s3.amazonaws.com/h2o-public-test-data/smalldata/junit/cars_20mpg.csv")
# convert response column to a factor
cars["economy_20mpg"] <- as.factor(cars["economy_20mpg"])
# set the predictor names and the response column name
predictors <- c("displacement","power","weight","acceleration","year")
response <- "economy_20mpg"
# split into train and validation
cars.splits <- h2o.splitFrame(data = cars, ratios = .8)
train <- cars.splits[[1]]
valid <- cars.splits[[2]]
# try using the `family` parameter:
cor_glm <- h2o.glm(x = predictors, y = response, family = 'binomial', training_frame = train,
validation_frame = valid)
# print the auc for your validation data
print(h2o.auc(cor_glm, valid = TRUE))
```

Hi Cluster,  
Can you please get me ...

# Client Only Passes Requests

Big Data Never Flows Through The Client  
Unless **Explicitly** Asked



The screenshot shows an RStudio interface with an 'Untitled1' script file open. The code performs the following steps:

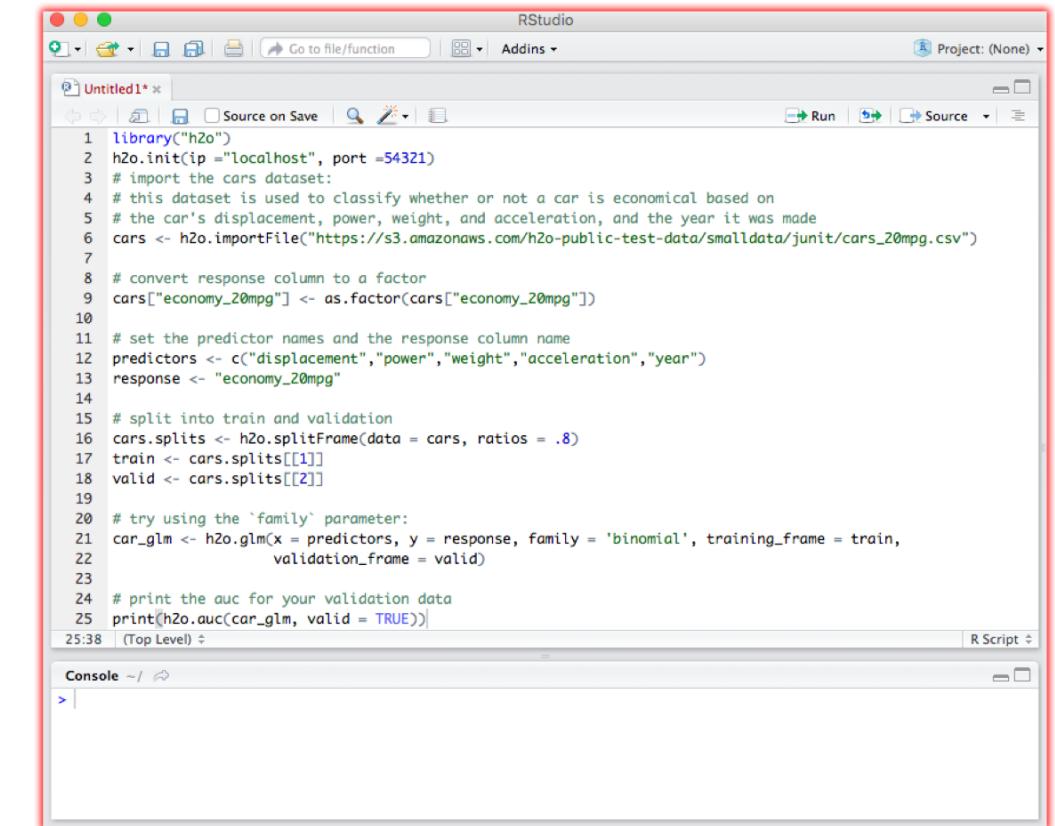
- Imports the 'h2o' library.
- Initializes H2O with port 54321.
- Imports the 'cars' dataset.
- Specifies the dataset is used to classify whether or not a car is economical based on its displacement, power, weight, and acceleration, and the year it was made.
- Imports the 'cars\_20mpg.csv' file from S3.
- Converts the 'economy\_20mpg' column to a factor.
- Sets predictor names and response column name.
- Creates splits for training and validation data.
- Attempts to use the 'family' parameter with 'h2o.glm'.
- Prints the AUC for the validation data.

No, you take care of the  
heavy lifting

# What if?

## Pulling Big Data into R Can Overwhelm Your Session

**as.data.frame(my\_big\_dataframe)**



A screenshot of an RStudio interface. The top bar shows 'RStudio' and 'Untitled1\*'. The main area is a code editor with the following R code:

```
1 library("h2o")
2 h2o.init(ip = "localhost", port = 54321)
3 # import the cars dataset:
4 # this dataset is used to classify whether or not a car is economical based on
5 # the car's displacement, power, weight, and acceleration, and the year it was made
6 cars <- h2o.importFile("https://s3.amazonaws.com/h2o-public-test-data/smalldata/junit/cars_20mpg.csv")
7
8 # convert response column to a factor
9 cars["economy_20mpg"] <- as.factor(cars["economy_20mpg"])
10
11 # set the predictor names and the response column name
12 predictors <- c("displacement","power","weight","acceleration","year")
13 response <- "economy_20mpg"
14
15 # split into train and validation
16 cars.splits <- h2o.splitFrame(data = cars, ratios = .8)
17 train <- cars.splits[[1]]
18 valid <- cars.splits[[2]]
19
20 # try using the `family` parameter:
21 car_glm <- h2o.glm(x = predictors, y = response, family = 'binomial', training_frame = train,
22                      validation_frame = valid)
23
24 # print the auc for your validation data
25 print(h2o.auc(car_glm, valid = TRUE))
```

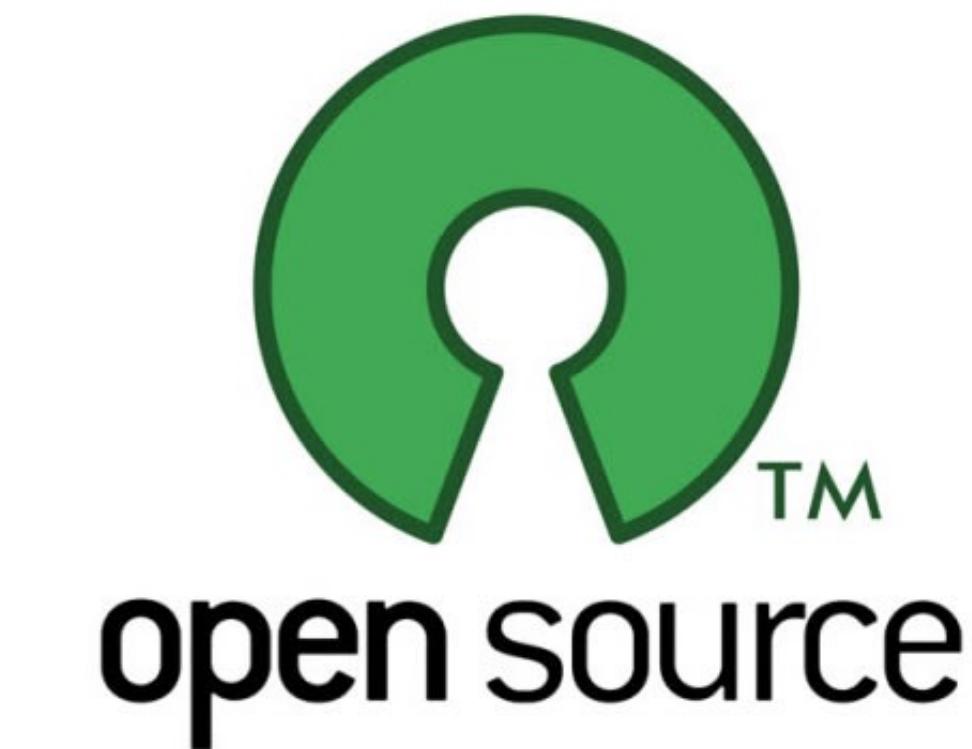
The bottom panel is the 'Console' window, which is currently empty.

AH!

# H2O Resources

- Documentation: <http://docs.h2o.ai>
- Tutorials: <https://github.com/h2oai/h2o-tutorials>
- Slidedecks: <https://github.com/h2oai/h2o-meetups>
- Videos: <https://www.youtube.com/user/0xdata>
- Stack Overflow: <https://stackoverflow.com/tags/h2o>
- Google Group: <https://tinyurl.com/h2ostream>
- Gitter: <http://gitter.im/h2oai/h2o-3>
- Events & Meetups: <http://h2o.ai/events>

# Contribute to H2O!



Get in touch over email, Gitter or JIRA.

<https://github.com/h2oai/h2o-3/blob/master/CONTRIBUTING.md>

# Thank you!

@navdeep-G on Github

@Navdeep\_Gill\_ on Twitter

navdeep@h2o.ai